

## nvae-on-mnist

September 25, 2021

```
[ ]: !nvidia-smi
```

Sat Sep 25 10:01:55 2021

```
+-----+  
| NVIDIA-SMI 470.63.01      Driver Version: 460.32.03      CUDA Version: 11.2      |  
+-----+-----+-----+  
| GPU  Name      Persistence-M| Bus-Id      Disp.A  | Volatile Uncorr. ECC  | | | | |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M.  |  
|          |          |          |          |          |          |          MIG M.  |  
+=====+=====+=====+=====+=====+=====+=====+  
|  0  Tesla K80          Off  | 00000000:00:04.0 Off  |                  0  | | | |
| N/A    73C     P8    34W / 149W |      0MiB / 11441MiB |      0%  Default  |  
|          |          |          |          |          |          N/A  |  
+-----+-----+-----+-----+-----+-----+  
  
+-----+  
| Processes:  
| GPU  GI  CI      PID  Type  Process name          GPU Memory  |  
|          ID  ID            |          |          |          Usage  |  
+=====+=====+=====+=====+=====+=====+=====+  
|  No running processes found  
+-----+
```

```
[ ]: import cv2  
import os  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import torch  
import torchvision  
import torch.nn as nn  
from torchvision import transforms  
import torch.nn.functional as F  
from torchvision.utils import save_image
```

```
[ ]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
[ ]: device
```

```
[ ]: device(type='cuda')
```

## 0.1 Depth Wise Separable Convolution

This convolution originated from the idea that depth and spatial dimension of a filter can be separated- thus the name separable. You can separate the height and width dimension of these filters. Gx filter (see fig 3) can be viewed as matrix product of [1 2 1] transpose with [-1 0 1].

that the filter had disguised itself. It shows it had 9 parameters but it has actually 6. This has been possible because of separation of its height and width dimensions. The same idea applied to separate depth dimension from horizontal (widthheight) *gives us depth-wise separable convolution where we perform depth-wise convolution and after that we use a 11 filter to cover the depth dimension.*

### 0.1.1 Torch Implementation Related Document

```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,  
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

Parameters:

stride controls the stride for the cross-correlation, a single number or a tuple.

padding controls the amount of padding applied to the input. It can be either a string {'valid'}

dilation controls the spacing between the kernel points; also known as the à trous algorithm. [

groups controls the connections between inputs and outputs. in\_channels and out\_channels must be

```
[ ]: class depthwise_separable_conv(nn.Module):  
    def __init__(self, nin, kernels_per_layer, nout):  
        super(depthwise_separable_conv, self).__init__()  
        self.depthwise = nn.Conv2d(nin, nin * kernels_per_layer, kernel_size=5, u  
        ↪padding=2, groups=nin)  
        self.pointwise = nn.Conv2d(nin * kernels_per_layer, nout, kernel_size=1)  
  
    def forward(self, x):  
        out = self.depthwise(x)  
        out = self.pointwise(out)  
        return out
```

## 0.2 Swish Activation Function

Swish is a smooth, non-monotonic function that consistently matches or outperforms ReLU on deep networks applied to a variety of challenging domains such as Image classification and Machine translation. It is unbounded above and bounded below & it is the non-monotonic attribute that actually creates the difference. With self-gating, it requires just a scalar input whereas in multi-gating scenario, it would require multiple two-scalar input.

```
[ ]: def swish(x):
    return x * torch.sigmoid(x)
```

### 0.3 Squeeze and Excitation Networks

For more information see this paper.

Squeeze-and-Excitation Networks (SENets) introduce a building block for CNNs that improves channel interdependencies at almost no computational cost.

The transformation simply corresponds with the operation that the network where you are going to implement the SE block would perform in its natural scheme. For instance, if you are in a block within a ResNet, the Ftr term will correspond with the process of the entire residual block (convolution, batch normalization, ReLU...).

The squeezing step is probably the most simply one. It basically performs a average pooling at each channel to create a 1x1 squeezed representation of the volume U.

The authors introduce a new parameter called the reduction ratio r, to introduce a first fully connected (FC) layer with a ReLU activation, before the gating network with the sigmoid activation.

The reason to do this is to introduce a bottleneck that allows us to reduce the dimensionality at the same time that introduce new non-linearities.

Furthermore, we can have better control on the model complexity and aid the generalization property of the network.

Having two FC layers will result on having 2 matrices of weights that will be learned by the network during the training in an end-to-end fashion (all of them are backpropagated together with the convolutional kernels).

The last step, scaling, is indeed a re-scaling operation. We are going to give the squeezed vector its original shape, keeping the information obtained during the excitation step.

Mathematically, the scaling is achieved by simple scalar product of each channel on the input volume with the corresponding channel on the activated 1x1 squeezed vector.

```
[ ]: class ChannelSELayer(nn.Module):
    def __init__(self, num_channels, reduction_ratio=2):
        super(ChannelSELayer, self).__init__()
        num_channels_reduced = num_channels // reduction_ratio
        self.reduction_ratio = reduction_ratio
        self.fc1 = nn.Linear(num_channels, num_channels_reduced, bias=True)
        self.fc2 = nn.Linear(num_channels_reduced, num_channels, bias=True)
        self.relu = nn.ReLU()
        self.sigmoid = nn.Sigmoid()

    def forward(self, input_tensor):
        batch_size, num_channels, H, W = input_tensor.size()
        squeeze_tensor = input_tensor.view(batch_size, num_channels, -1).
        ↵mean(dim=2)
```

```

        fc_out_1 = self.relu(self.fc1(squeeze_tensor))
        fc_out_2 = self.sigmoid(self.fc2(fc_out_1))

        a, b = squeeze_tensor.size()
        output_tensor = torch.mul(input_tensor, fc_out_2.view(a, b, 1, 1))
        return output_tensor

```

```

[ ]: class dec_res(nn.Module):
    def __init__(self,in_channel):
        super(dec_res,self).__init__()
        self.bn1 = nn.BatchNorm2d(in_channel)
        self.c1 = nn.
        →Conv2d(in_channels=in_channel,out_channels=2*in_channel,kernel_size=1,stride=1,padding=0)
        self.bn2 = nn.BatchNorm2d(2*in_channel)
        self.dc1 =
        →depthwise_separable_conv(nin=2*in_channel,kernels_per_layer=3,nout=2*in_channel)
        self.bn3 = nn.BatchNorm2d(2*in_channel)
        self.c2 = nn.
        →Conv2d(in_channels=2*in_channel,out_channels=in_channel,kernel_size=1,stride=1,padding=0)
        self.bn4 = nn.BatchNorm2d(in_channel)
        self.SE = ChannelSELayer(in_channel)
    def forward(self,x1):
        x = self.c1(self.bn1(x1))
        x = swish(self.bn2(x))
        x = self.dc1(x)
        x = swish(self.bn3(x))
        x = self.bn4(self.c2(x))
        x = self.SE(x)
        return x+x1

```

```

[ ]: class enc_res(nn.Module):
    def __init__(self,in_channel):
        super(enc_res,self).__init__()
        self.bn1 = nn.BatchNorm2d(in_channel)
        self.c1 = nn.
        →Conv2d(in_channels=in_channel,out_channels=2*in_channel,kernel_size=3,stride=1,padding=1)
        self.bn2 = nn.BatchNorm2d(2*in_channel)
        self.c2 = nn.
        →Conv2d(in_channels=2*in_channel,out_channels=in_channel,kernel_size=3,stride=1,padding=1)
        self.bn3 = nn.BatchNorm2d(in_channel)
        self.SE = ChannelSELayer(in_channel)
    def forward(self,x1):
        x = self.c1(swish(self.bn1(x1)))
        x = self.c2(swish(self.bn2(x)))
        x = self.SE(x)
        return x+x1

```

```
[ ]: class NVAE(nn.Module):
    def __init__(self,start_channel,original_dim):
        super(NVAE,self).__init__()
        self.original_dim = original_dim
        self.conv1 = nn.
        ↪Conv2d(in_channels=start_channel,out_channels=8,kernel_size=3,stride=1,padding=1)
        self.encblock1 = enc_res(8)
        self.dsconv1 = nn.
        ↪Conv2d(in_channels=8,out_channels=8,kernel_size=2,stride=2,padding=0)
        self.encblock2 = enc_res(8)
        self.dsconv2 = nn.
        ↪Conv2d(in_channels=8,out_channels=8,kernel_size=2,stride=2,padding=0)

        self.qmu1 = nn.
        ↪Linear(original_dim*original_dim*2,original_dim*original_dim*2)
        self.qvar1 = nn.
        ↪Linear(original_dim*original_dim*2,original_dim*original_dim*2)

        self.qmu0 = nn.Linear(original_dim*original_dim//
        ↪2,original_dim*original_dim//2)
        self.qvar0 = nn.Linear(original_dim*original_dim//
        ↪2,original_dim*original_dim//2)

        self.pmu1 = nn.
        ↪Linear(original_dim*original_dim*2,original_dim*original_dim*2)
        self.pvar1 = nn.
        ↪Linear(original_dim*original_dim*2,original_dim*original_dim*2)

        self.decblock1 = dec_res(8)
        self.usconv1 = nn.
        ↪ConvTranspose2d(in_channels=8,out_channels=8,kernel_size=2,stride=2,padding=0)
        self.decblock2 = dec_res(16)
        self.usconv2 = nn.
        ↪ConvTranspose2d(in_channels=16,out_channels=16,kernel_size=2,stride=2,padding=0)
        self.decblock3 = dec_res(16)
        self.finconv = nn.
        ↪Conv2d(in_channels=16,out_channels=start_channel,kernel_size=3,stride=1,padding=1)

    def forward(self,x):
        z1 = self.dsconv1(self.encblock1(self.conv1(x)))
        z0 = self.dsconv2(self.encblock2(z1))

        qmu0 = self.qmu0(z0.reshape(z0.shape[0],self.original_dim*self.original_dim/
        ↪2))
        qvar0 = self.qvar0(z0.reshape(z0.shape[0],self.original_dim*self.
        ↪original_dim//2))
```

```

qmu1 = self.qmu1(z1.reshape(z1.shape[0],self.original_dim*self.
˓→original_dim*2))
qvar1 = self.qvar1(z1.reshape(z1.shape[0],self.original_dim*self.
˓→original_dim*2))

stdvar0 = qvar0.mul(0.5).exp_()
stdvar1 = qvar1.mul(0.5).exp_()

e0 = torch.randn(qmu0.shape).to(device)
ez0 = qmu0+e0*stdvar0
ez0 = ez0.reshape(ez0.shape[0],8,self.original_dim//4,self.original_dim//4)
ez1 = self.usconv1(self.decblock1(ez0))

pmu1 = self.pmu1(ez1.reshape(ez1.shape[0],self.original_dim*self.
˓→original_dim*2))
pvar1 = self.pvar1(ez1.reshape(ez1.shape[0],self.original_dim*self.
˓→original_dim*2))

pstdvar1 = pvar1.mul(0.5).exp_()

e2 = torch.randn(qmu1.shape).to(device)
ez2 = pmu1+qmu1 + e2*pstdvar1*stdvar1
ez2 = ez2.reshape(ez2.shape[0],8,self.original_dim//2,self.original_dim//2)

final = torch.cat((ez1,ez2),1)

recons = nn.Sigmoid()(self.finconv(self.decblock3(self.usconv2(self.
˓→decblock2(final)))))

return qmu0,qvar0,qmu1,qvar1,pmu1,pvar1,recons

def sample(self,bs):
    e = torch.randn([bs,8,self.original_dim//4,self.original_dim//4]).to(device)
    ez1 = self.usconv1(self.decblock1(e))

    pmu1 = self.pmu1(ez1.reshape(ez1.shape[0],self.original_dim*self.
˓→original_dim*2))
    pvar1 = self.pvar1(ez1.reshape(ez1.shape[0],self.original_dim*self.
˓→original_dim*2))

    stdvar1 = pvar1.mul(0.5).exp_()

    e1 = torch.randn([ez1.shape[0],self.original_dim*self.original_dim*2]).to(device)
    e1 = pmu1 + e1*stdvar1

```

```

e1 = e1.reshape(e1.shape[0],8,self.original_dim//2,self.original_dim//2)
recons = nn.Sigmoid()(self.finconv(self.decblock3(self.usconv2(self.
    ↳decblock2(torch.cat((ez1,e1),1))))))

return recons

def loss(self,x):
    qmu0,qvar0,qmu1,qvar1,pmu1,pvar1,recons = self.forward(x)
    klz0 = 0.5*torch.sum(torch.square(qmu0)+qvar0.exp()-qvar0-1)/x.shape[0]
    klz1 = 0.5*torch.sum(torch.square(qmu1)/pvar1.exp()+qvar1.exp()-qvar1-1)
    reconsloss = nn.BCELoss()(recons,x)
    return klz0,klz1,reconsloss

```

[ ]: batch\_size=64

[ ]: transform = transforms.Compose([transforms.ToTensor()])

```

[ ]: train_dataset = torchvision.datasets.MNIST(root='data/mnist',
                                                train=True,
                                                transform=transform,
                                                download=True)

test_dataset = torchvision.datasets.MNIST(root='data/mnist',
                                           train=False,
                                           transform=transform)

train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
                                            batch_size=batch_size,
                                            shuffle=True)

test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
                                           batch_size=batch_size,
                                           shuffle=True)

```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz  
 Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to  
 data/mnist/MNIST/raw/train-images-idx3-ubyte.gz

0%| 0/9912422 [00:00<?, ?it/s]

Extracting data/mnist/MNIST/raw/train-images-idx3-ubyte.gz to  
 data/mnist/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz  
 Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to  
 data/mnist/MNIST/raw/train-labels-idx1-ubyte.gz

0%| 0/28881 [00:00<?, ?it/s]

```

Extracting data/mnist/MNIST/raw/train-labels-idx1-ubyte.gz to
data/mnist/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to
data/mnist/MNIST/raw/t10k-images-idx3-ubyte.gz

0%|          | 0/1648877 [00:00<?, ?it/s]

Extracting data/mnist/MNIST/raw/t10k-images-idx3-ubyte.gz to
data/mnist/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to
data/mnist/MNIST/raw/t10k-labels-idx1-ubyte.gz

0%|          | 0/4542 [00:00<?, ?it/s]

Extracting data/mnist/MNIST/raw/t10k-labels-idx1-ubyte.gz to
data/mnist/MNIST/raw

```

```

/usr/local/lib/python3.7/dist-packages/torchvision/datasets/mnist.py:498:
UserWarning: The given NumPy array is not writeable, and PyTorch does not
support non-writeable tensors. This means you can write to the underlying
(supposedly non-writeable) NumPy array using the tensor. You may want to copy
the array to protect its data or make it writeable before converting it to a
tensor. This type of warning will be suppressed for the rest of this program.
(Triggered internally at  /pytorch/torch/csrc/utils/tensor_numpy.cpp:180.)
    return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)

```

```
[ ]: model = NVAE(1,28).to(device)
```

AdaMax is a generalisation of Adam from l<sub>2</sub> the norm l<sub>∞</sub> to the norm.

```
torch.optim.Adamax(params, lr=0.002, betas=(0.9, 0.999), eps=1e-08, weight_decay=0)
```

Implements Adamax algorithm (a variant of Adam based on infinity norm).

```
[ ]: optim = torch.optim.Adamax(model.parameters())
```

```
[ ]: epochs=50
```

```
[ ]: import matplotlib.pyplot as plt
from mpl_toolkits.axes_grid1 import ImageGrid
```

```
[ ]: for epoch in range(epochs):
    minloss = 1
    running_k10_loss=0
    running_recons_loss=0
    running_k11_loss=0
```

```

num_images=0
for i,(img,label) in enumerate(train_loader):
    img = img.to(device)
    optim.zero_grad()
    klz0,klz1,recons = model.loss(img)
    loss=recons+epoch*0.0001*klz0+epoch*0.0001*klz1
    loss.backward()
    optim.step()
    running_kl0_loss = running_kl0_loss + klz0.item()*len(img)
    running_kl1_loss = running_kl1_loss + klz1.item()*len(img)
    running_recons_loss = running_recons_loss + recons.item()*len(img)

    num_images= num_images+len(img)
    print('epoch: '+str(epoch)+ ' kl0_loss: '+str(running_kl0_loss/num_images)+ ' ↵
    ↵recons_loss: '+str(running_recons_loss/num_images)+ ' kl1_loss: ↵
    ↵'+str(running_kl1_loss/num_images))
    imgs = model.sample(64).cpu().detach().reshape(64,28,28)
    plt.gray()
    fig = plt.figure(figsize=(8., 8.))
    grid = ImageGrid(fig, 111, nrows_ncols=(8, 8), axes_pad=0.05)

    for ax, im in zip(grid, imgs):
        ax.imshow(im)
    plt.savefig(str(epoch)+".png")

```

epoch: 0 kl0\_loss: 751.782080317688 recons\_loss: 0.08109162319699924 kl1\_loss:  
8917296.777019141  
epoch: 1 kl0\_loss: 570.4323500651042 recons\_loss: 0.15105966951052346 kl1\_loss:  
38005.96190992839  
epoch: 2 kl0\_loss: 402.2469321289062 recons\_loss: 0.14689265320301056 kl1\_loss:  
259.6337049641927  
epoch: 3 kl0\_loss: 213.27998752441405 recons\_loss: 0.1471416516462962 kl1\_loss:  
115.2509824239095  
epoch: 4 kl0\_loss: 115.06704755045573 recons\_loss: 0.1378506362915039 kl1\_loss:  
52.97208299967448  
epoch: 5 kl0\_loss: 73.83459227701823 recons\_loss: 0.12621081261634826 kl1\_loss:  
22.230131729888917  
epoch: 6 kl0\_loss: 57.30363482259114 recons\_loss: 0.1184653028011322 kl1\_loss:  
9.7949793665556804  
epoch: 7 kl0\_loss: 48.309914426676436 recons\_loss: 0.1144391380906105 kl1\_loss:  
4.822097951634725  
epoch: 8 kl0\_loss: 41.855485040283206 recons\_loss: 0.11107408939599991 kl1\_loss:  
2.3236456545511883  
epoch: 9 kl0\_loss: 36.97051285603841 recons\_loss: 0.1051474888086319 kl1\_loss:  
1.2710862868467967  
epoch: 10 kl0\_loss: 33.71670335693359 recons\_loss: 0.0998356648683548 kl1\_loss:  
0.6939344835599264

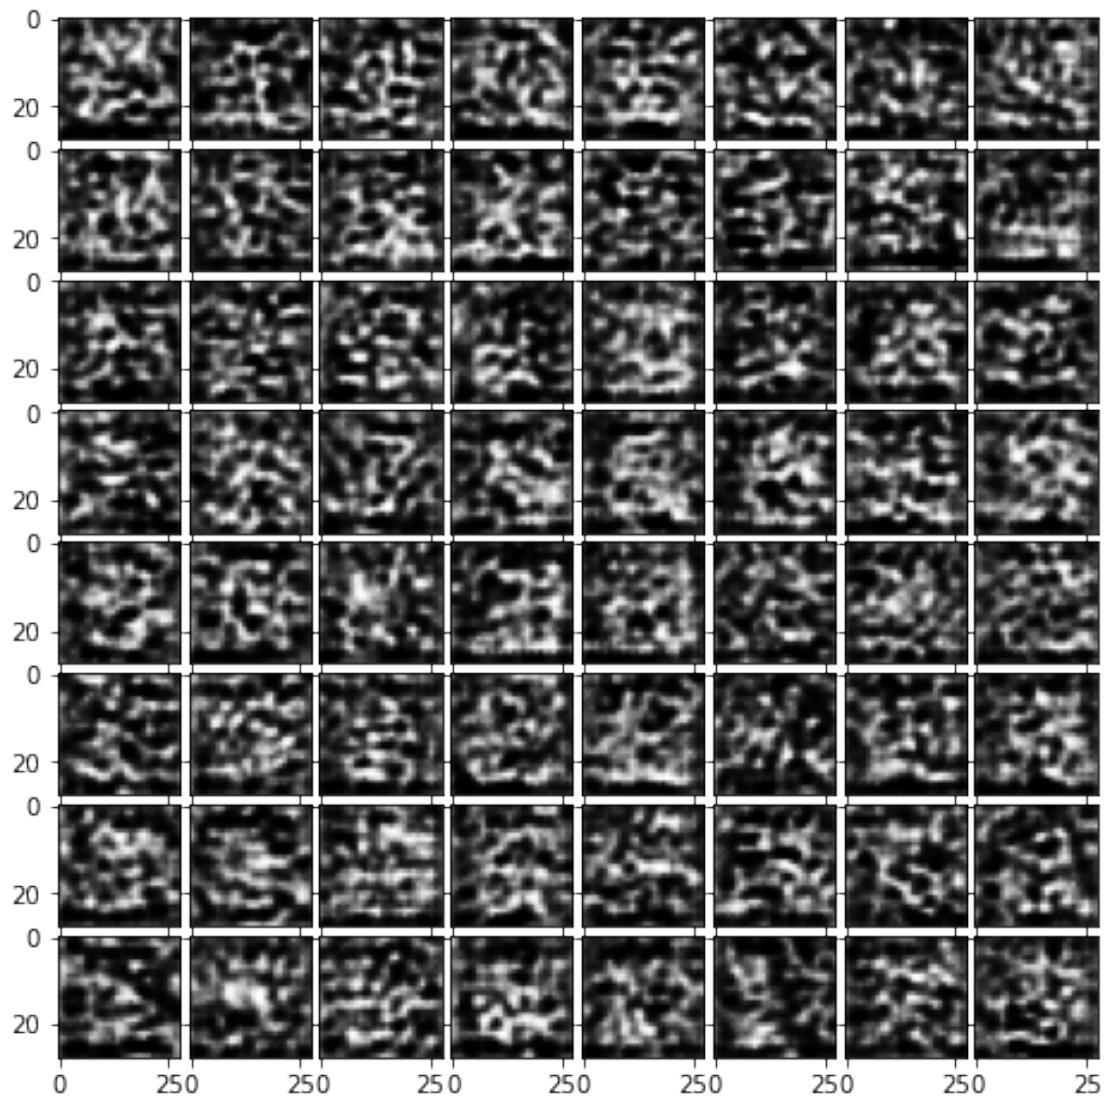
```
epoch: 11 k10_loss: 31.774424300130207 recons_loss: 0.09806553760369618
k11_loss: 0.4290548406124115
epoch: 12 k10_loss: 30.256083868408204 recons_loss: 0.09812449281613032
k11_loss: 0.2950733585993449
epoch: 13 k10_loss: 28.753330624389648 recons_loss: 0.09869587037960688
k11_loss: 0.2302874468008677
epoch: 14 k10_loss: 27.493141733805338 recons_loss: 0.09959638084967931
k11_loss: 0.18545938824017844
epoch: 15 k10_loss: 26.28852609049479 recons_loss: 0.10048087741136551 k11_loss:
0.15588229813575744
epoch: 16 k10_loss: 25.20113154602051 recons_loss: 0.1014794246951739 k11_loss:
0.13966897044181822
epoch: 17 k10_loss: 24.198837790934245 recons_loss: 0.10256501171588897
k11_loss: 0.12259357266426087
epoch: 18 k10_loss: 23.263466990152995 recons_loss: 0.10358741597731908
k11_loss: 0.10926611018180847
epoch: 19 k10_loss: 22.410162530517578 recons_loss: 0.10468868658542634
k11_loss: 0.10576434677441915

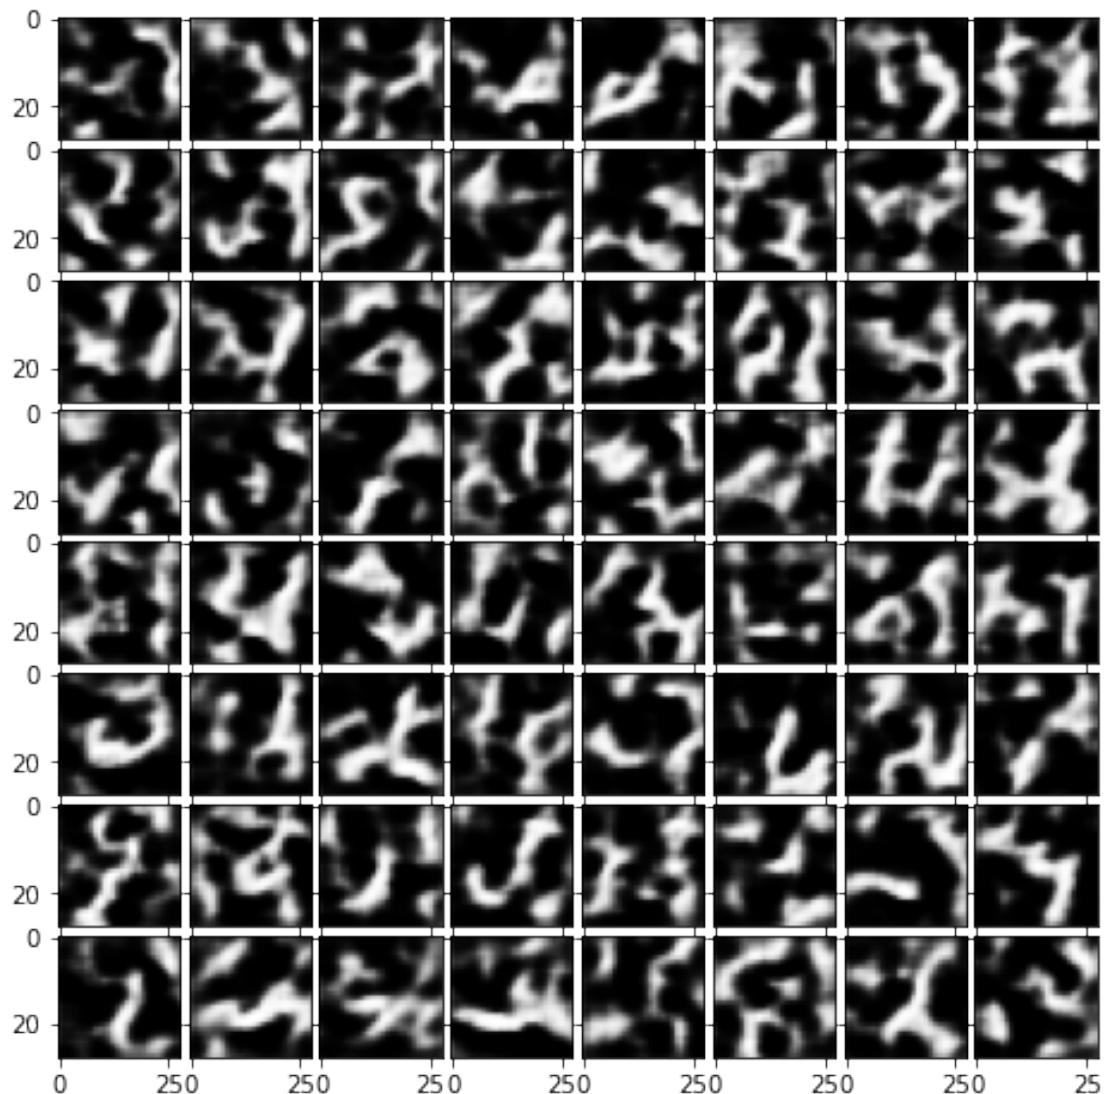
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:22: RuntimeWarning:
More than 20 figures have been opened. Figures created through the pyplot
interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and
may consume too much memory. (To control this warning, see the rcParam
`figure.max_open_warning`).

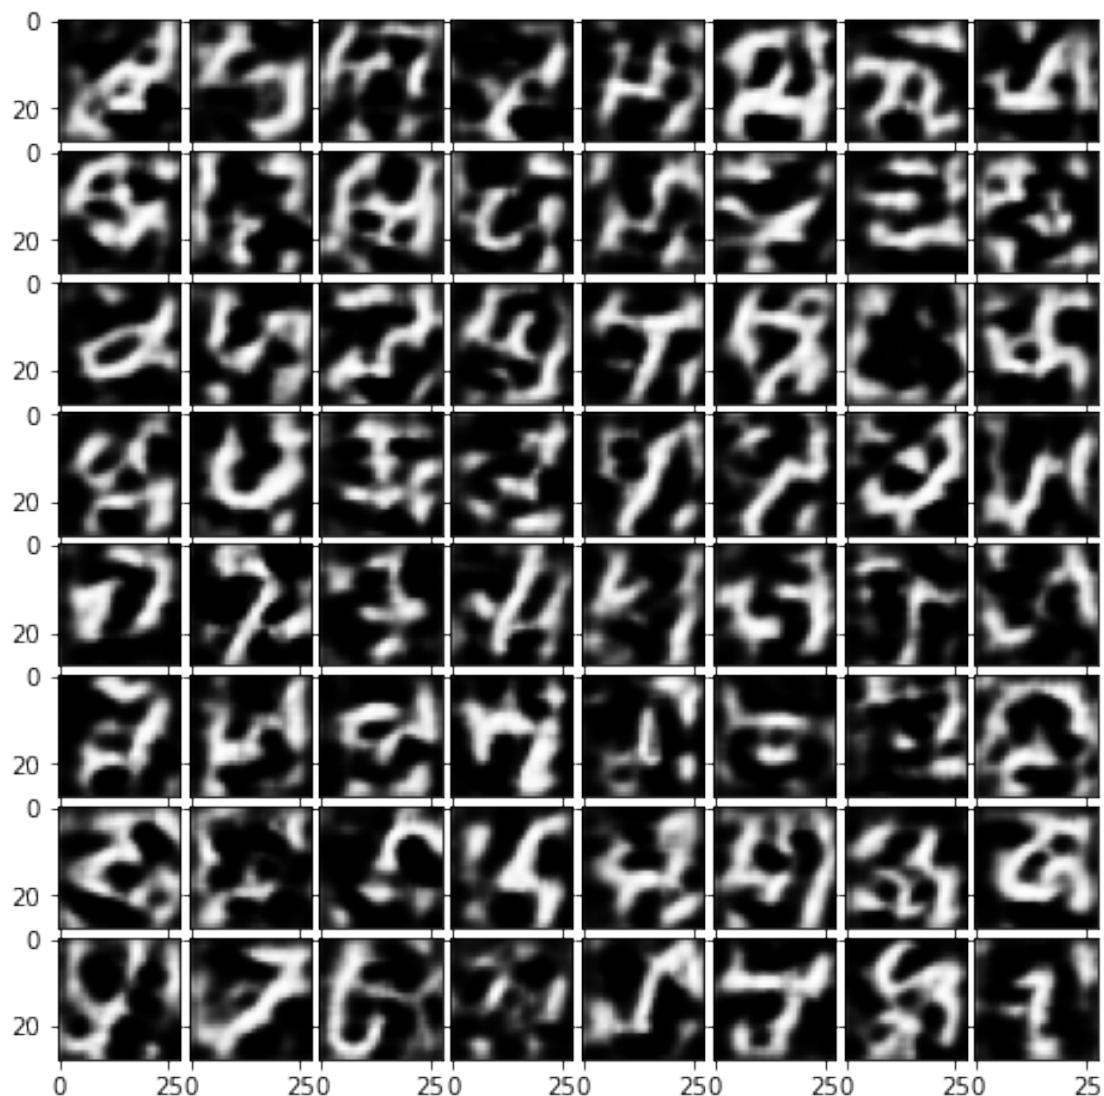
epoch: 20 k10_loss: 21.57954174499512 recons_loss: 0.1056950667778651 k11_loss:
0.09912984641393026
epoch: 21 k10_loss: 20.86861962381999 recons_loss: 0.10672866642077763 k11_loss:
0.09103862768809001
epoch: 22 k10_loss: 20.19998764851888 recons_loss: 0.10775813192129136 k11_loss:
0.08282333766619364
epoch: 23 k10_loss: 19.529254646809896 recons_loss: 0.10877485410372416
k11_loss: 0.07920589922269185
epoch: 24 k10_loss: 18.907623767089845 recons_loss: 0.10982415846586227
k11_loss: 0.07466760719617208
epoch: 25 k10_loss: 18.370249978637695 recons_loss: 0.1106940640091896 k11_loss:
0.07422551207542419
epoch: 26 k10_loss: 17.837274108886717 recons_loss: 0.11170441991885503
k11_loss: 0.0666663032690684
epoch: 27 k10_loss: 17.28653235066732 recons_loss: 0.11273801556825638 k11_loss:
0.06336531389554341
epoch: 28 k10_loss: 16.822976717122394 recons_loss: 0.11374140033721925
k11_loss: 0.06198535777727763
epoch: 29 k10_loss: 16.36267199198405 recons_loss: 0.11461928412914277 k11_loss:
0.06130962732632955
epoch: 30 k10_loss: 15.941362764485676 recons_loss: 0.11557901817162831
k11_loss: 0.058592124366760256
epoch: 31 k10_loss: 15.480578638203939 recons_loss: 0.11652140124638875
k11_loss: 0.05679292494455973
```

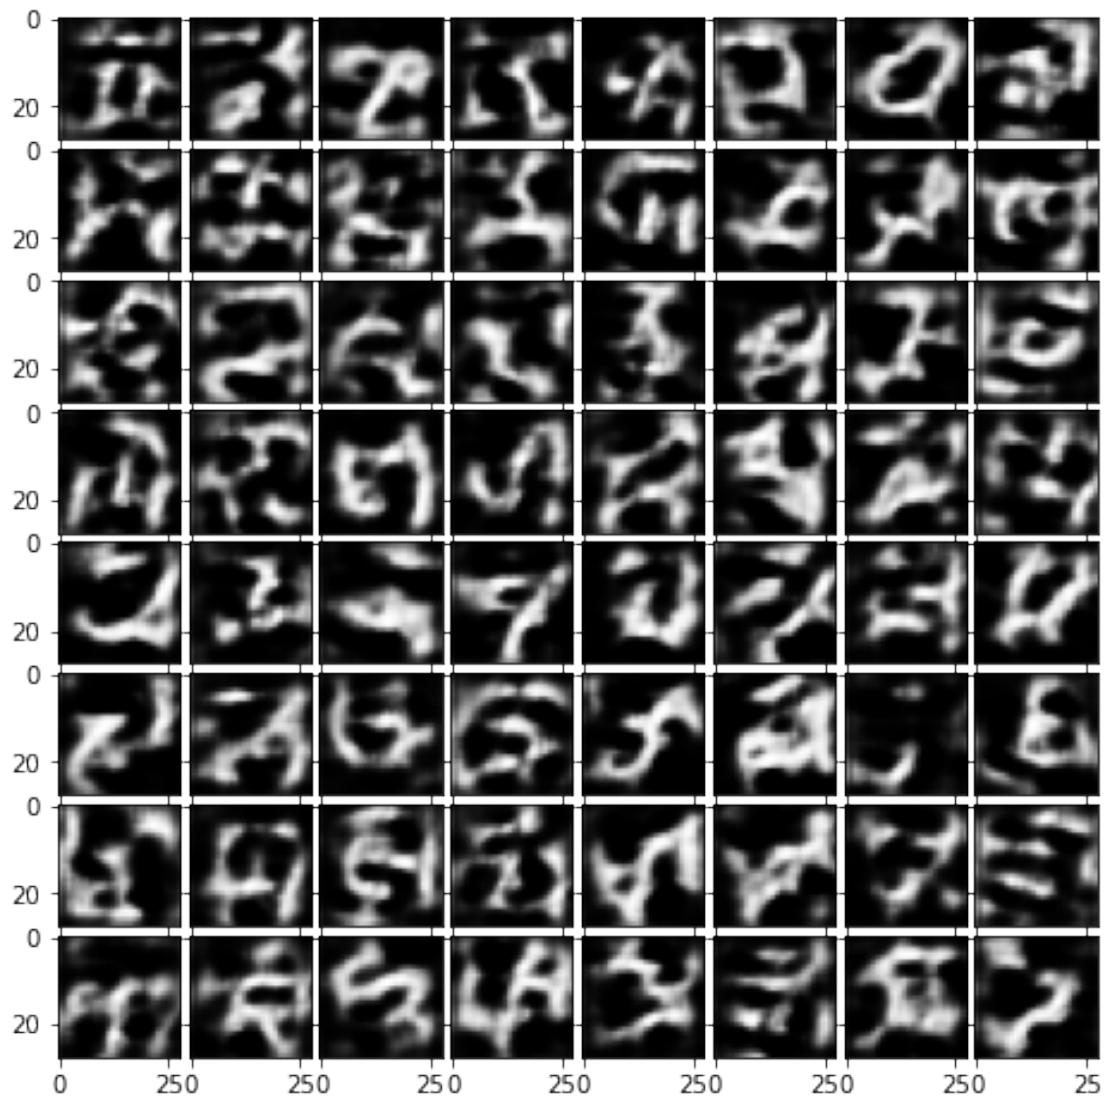
```
epoch: 32 k10_loss: 15.145828164164225 recons_loss: 0.11731142841180166
k11_loss: 0.05196991376876831
epoch: 33 k10_loss: 14.746272914632161 recons_loss: 0.11820075264771779
k11_loss: 0.05115602682431539
epoch: 34 k10_loss: 14.374892031351726 recons_loss: 0.11906812372207641
k11_loss: 0.04865165758132935
epoch: 35 k10_loss: 14.046451326497396 recons_loss: 0.11994835286537806
k11_loss: 0.04605357297261556
epoch: 36 k10_loss: 13.687309195963541 recons_loss: 0.12078918993473053
k11_loss: 0.044615216700236
epoch: 37 k10_loss: 13.387060451253255 recons_loss: 0.12152494059006373
k11_loss: 0.04352628908157349
epoch: 38 k10_loss: 13.072342435709636 recons_loss: 0.12241315124432246
k11_loss: 0.04461188589731852
epoch: 39 k10_loss: 12.81632827351888 recons_loss: 0.12308639443715413 k11_loss:
0.04126805640856425
epoch: 40 k10_loss: 12.538095556132 recons_loss: 0.12386690600713095 k11_loss:
0.03961378099123637
epoch: 41 k10_loss: 12.294520878092447 recons_loss: 0.12465567227204641
k11_loss: 0.04051355160077413
epoch: 42 k10_loss: 12.053678917948405 recons_loss: 0.12528542629877726
k11_loss: 0.038296764373779296
epoch: 43 k10_loss: 11.83693717549642 recons_loss: 0.12615564269224802 k11_loss:
0.03759724310239156
epoch: 44 k10_loss: 11.628213116455077 recons_loss: 0.1268004652897517 k11_loss:
0.0377593921661377
epoch: 45 k10_loss: 11.384304349263509 recons_loss: 0.12758560822804768
k11_loss: 0.03565698954264323
epoch: 46 k10_loss: 11.206871402994791 recons_loss: 0.12825051782131194
k11_loss: 0.03503577399253845
epoch: 47 k10_loss: 11.008865227254232 recons_loss: 0.12891203769048054
k11_loss: 0.03407250299453735
epoch: 48 k10_loss: 10.820389528401693 recons_loss: 0.12958042811552684
k11_loss: 0.032992552121480306
epoch: 49 k10_loss: 10.675146875 recons_loss: 0.13029568121433258 k11_loss:
0.03283892561594645

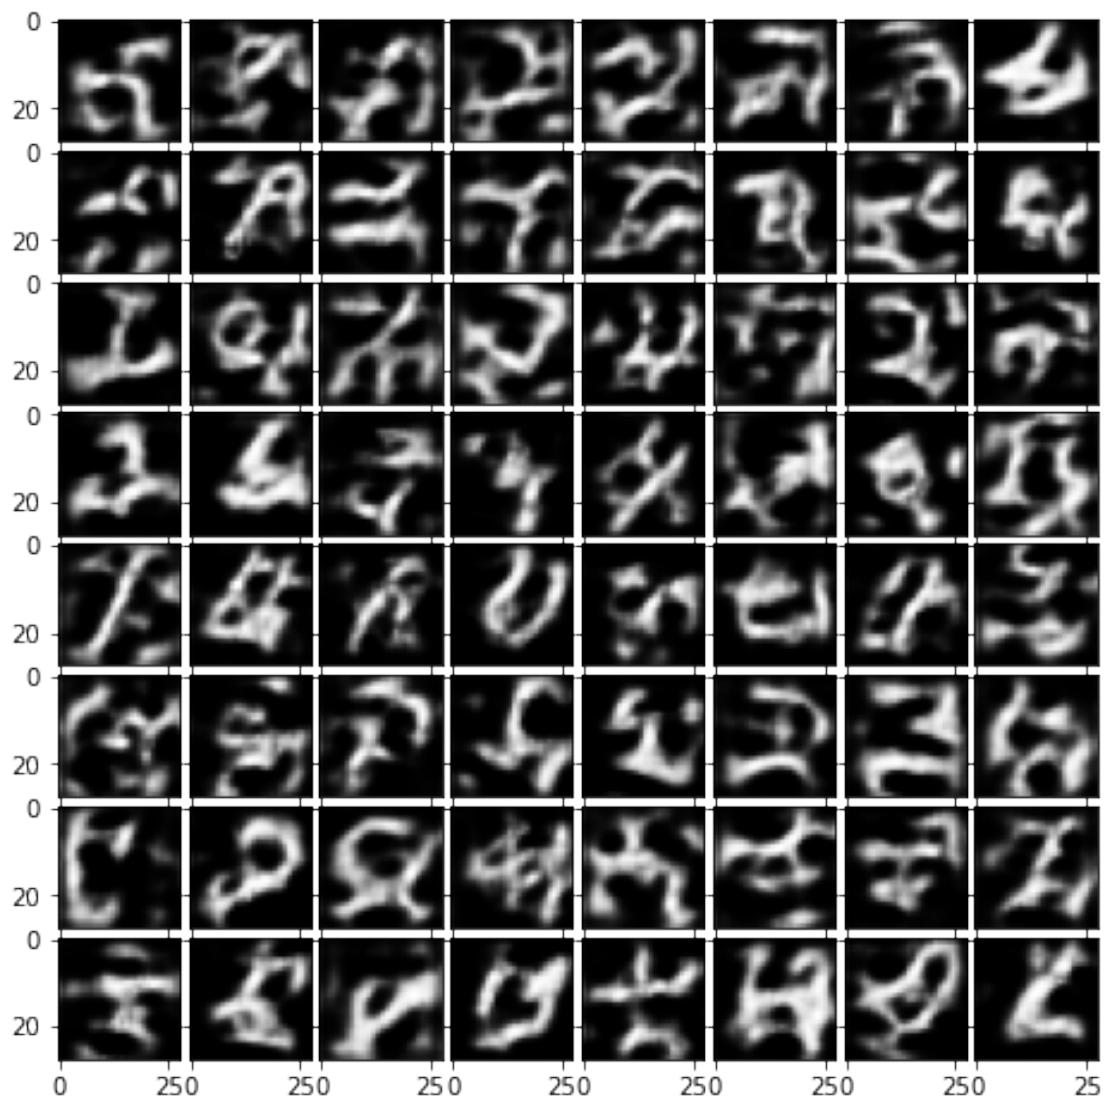
<Figure size 432x288 with 0 Axes>
```



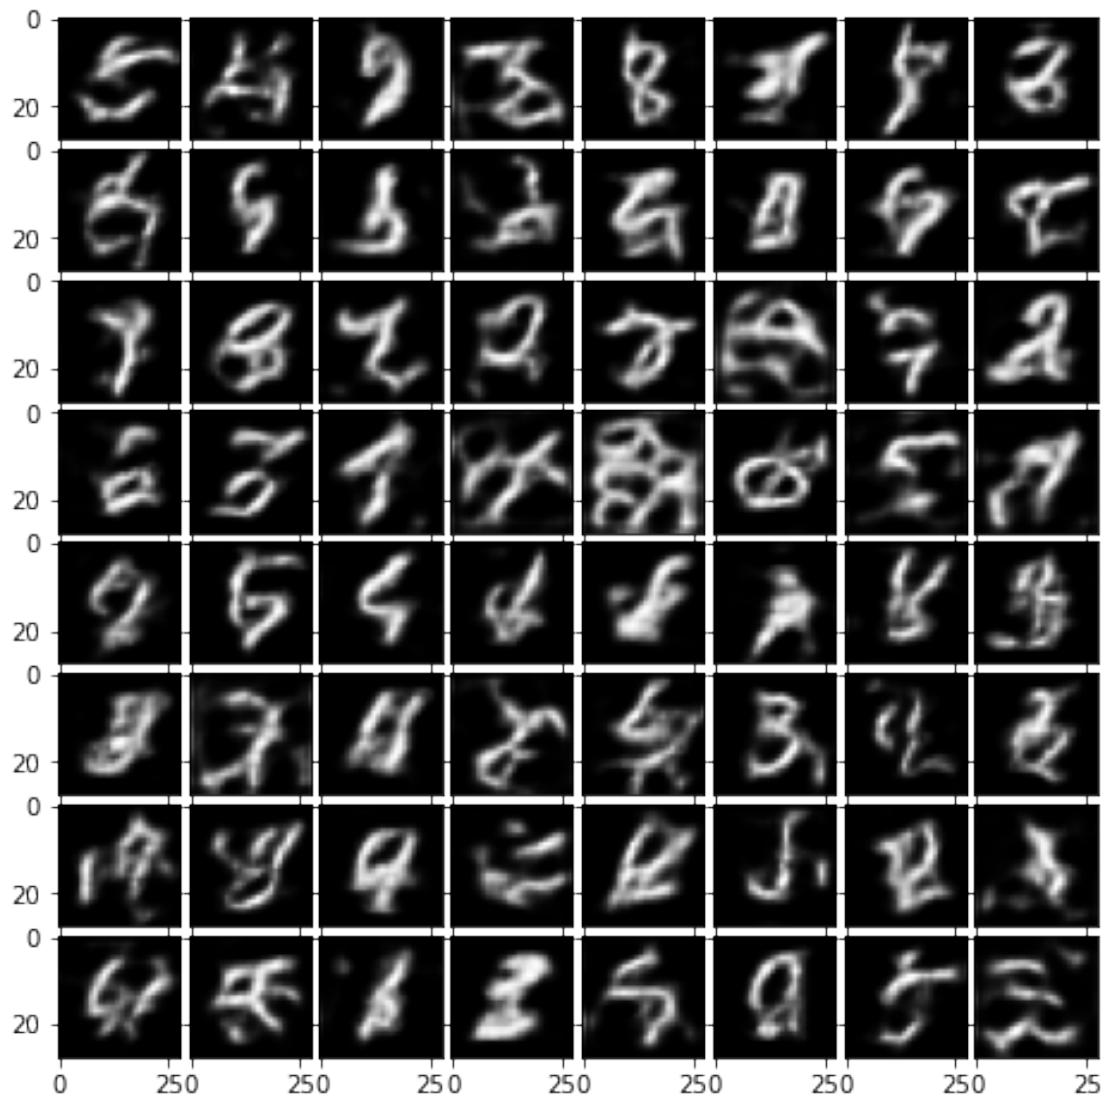




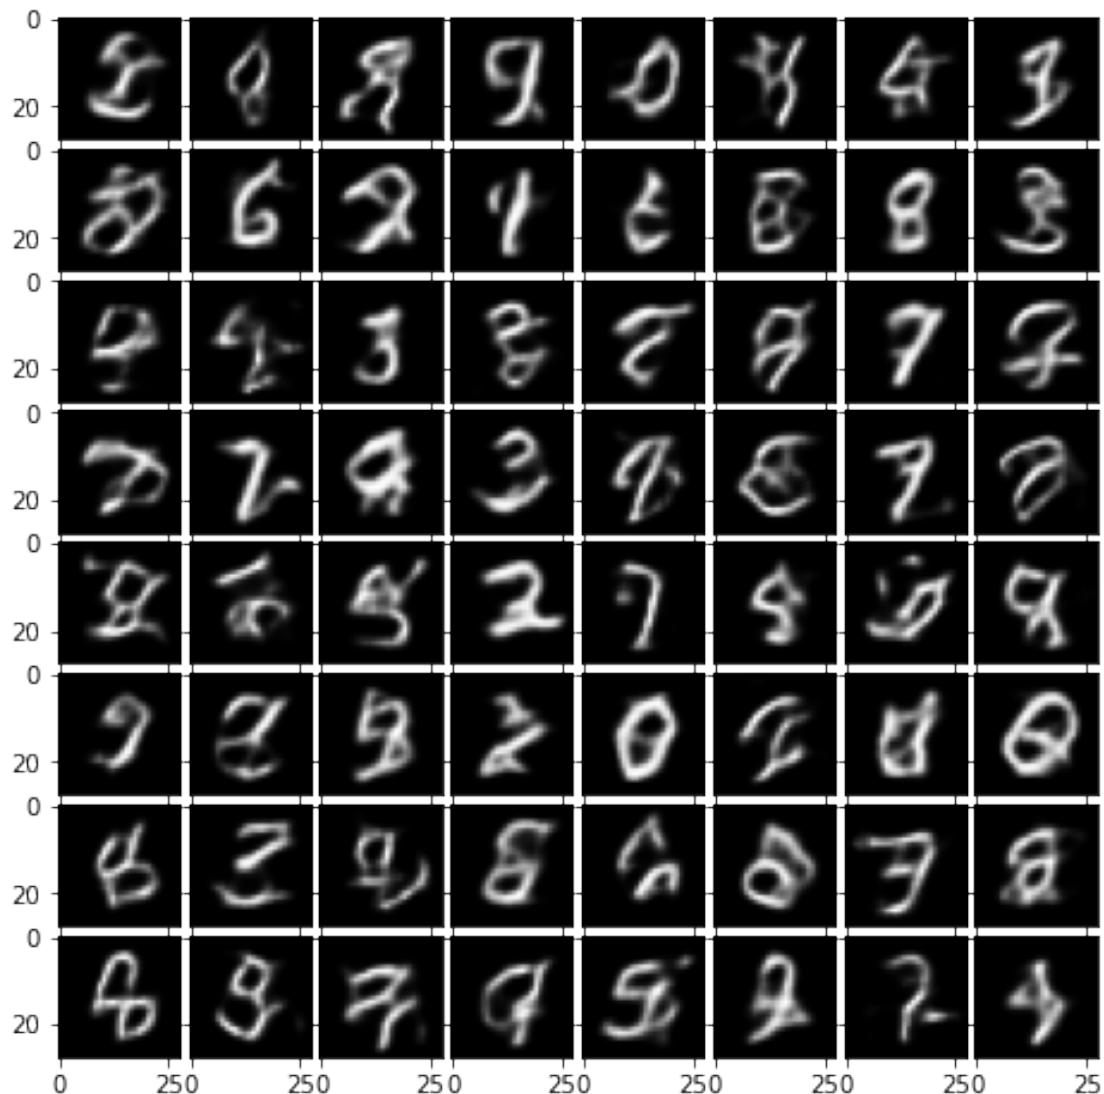


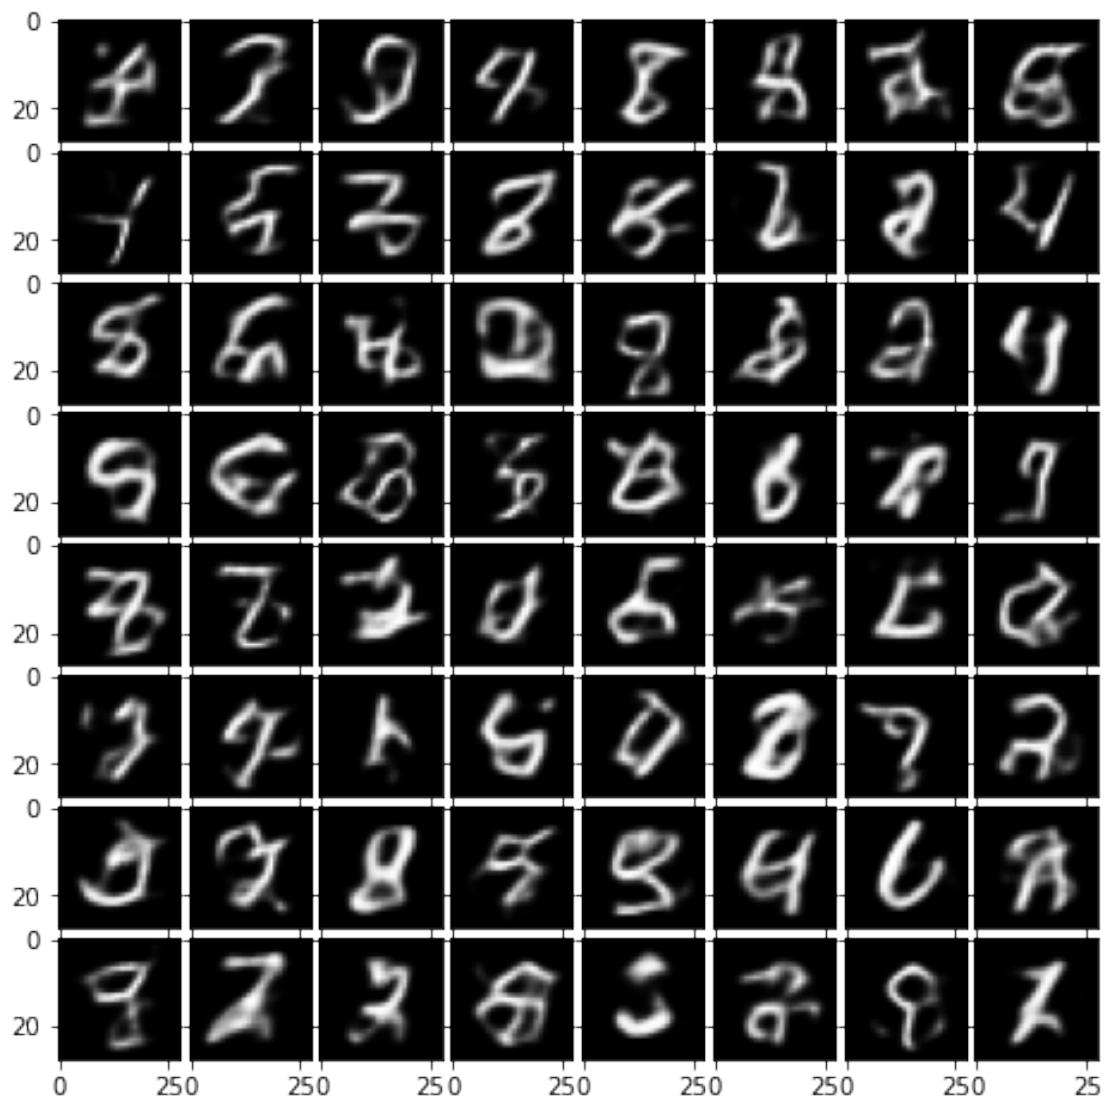




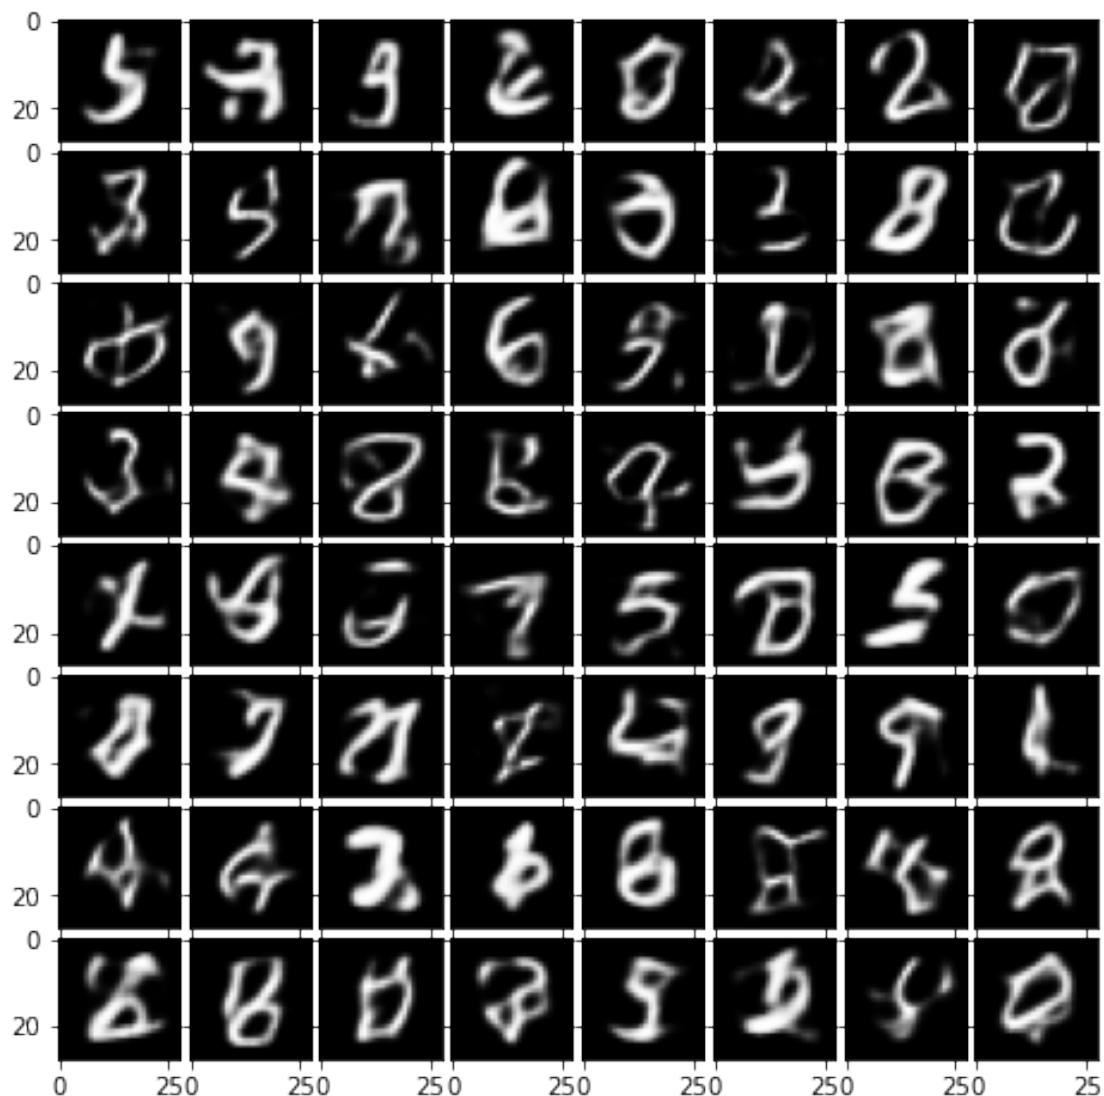


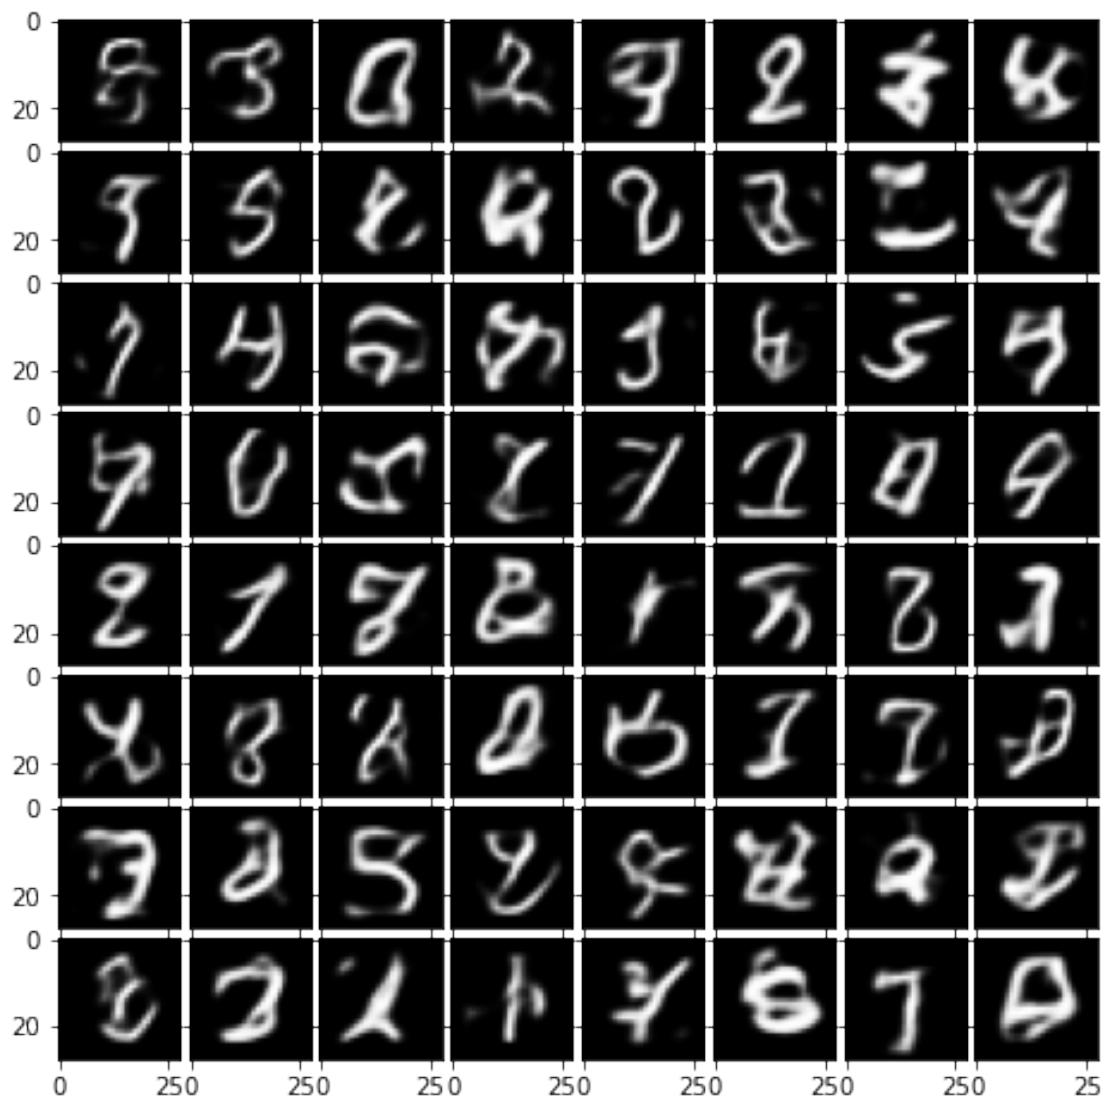


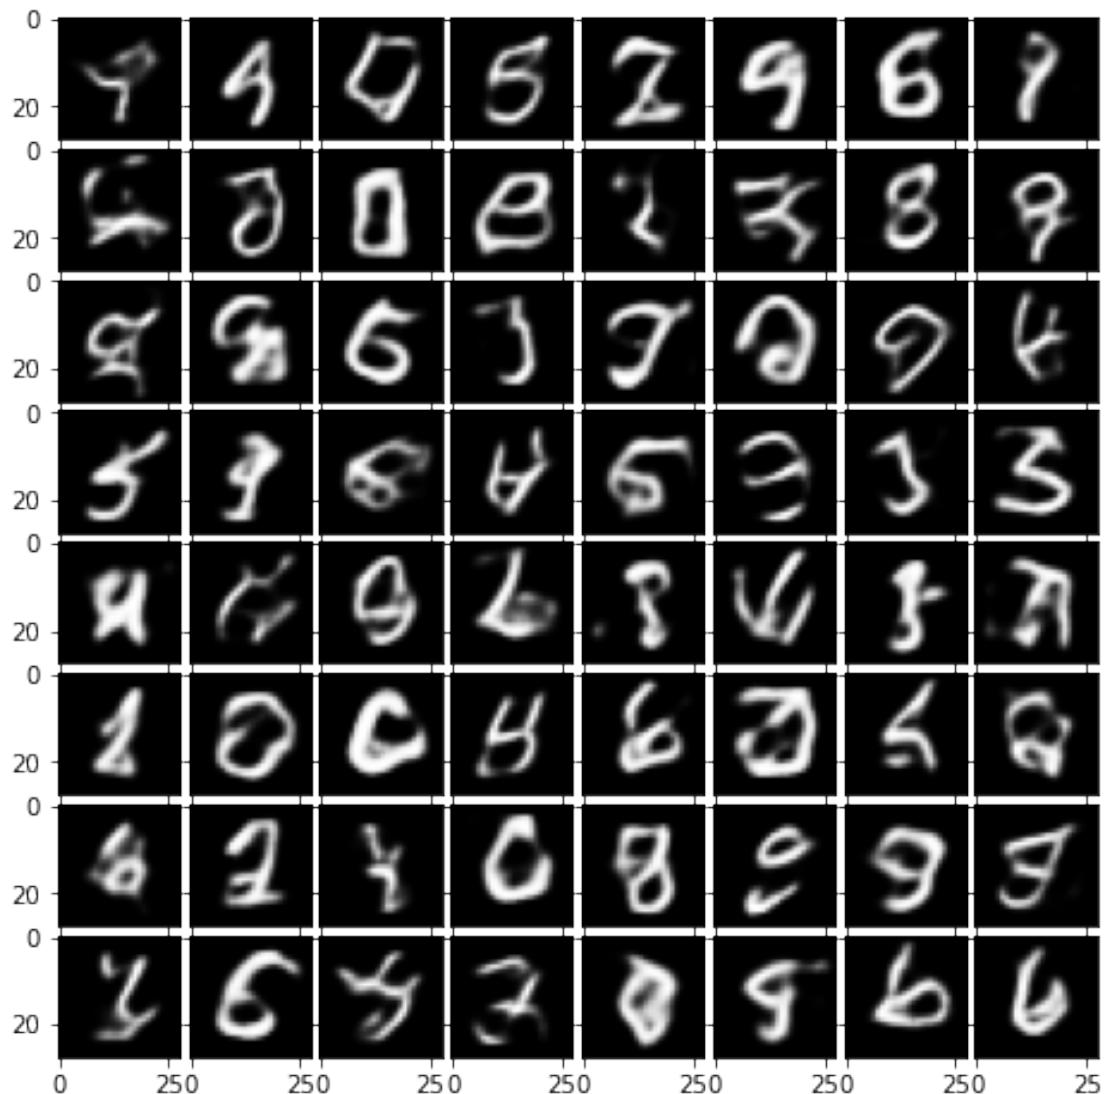


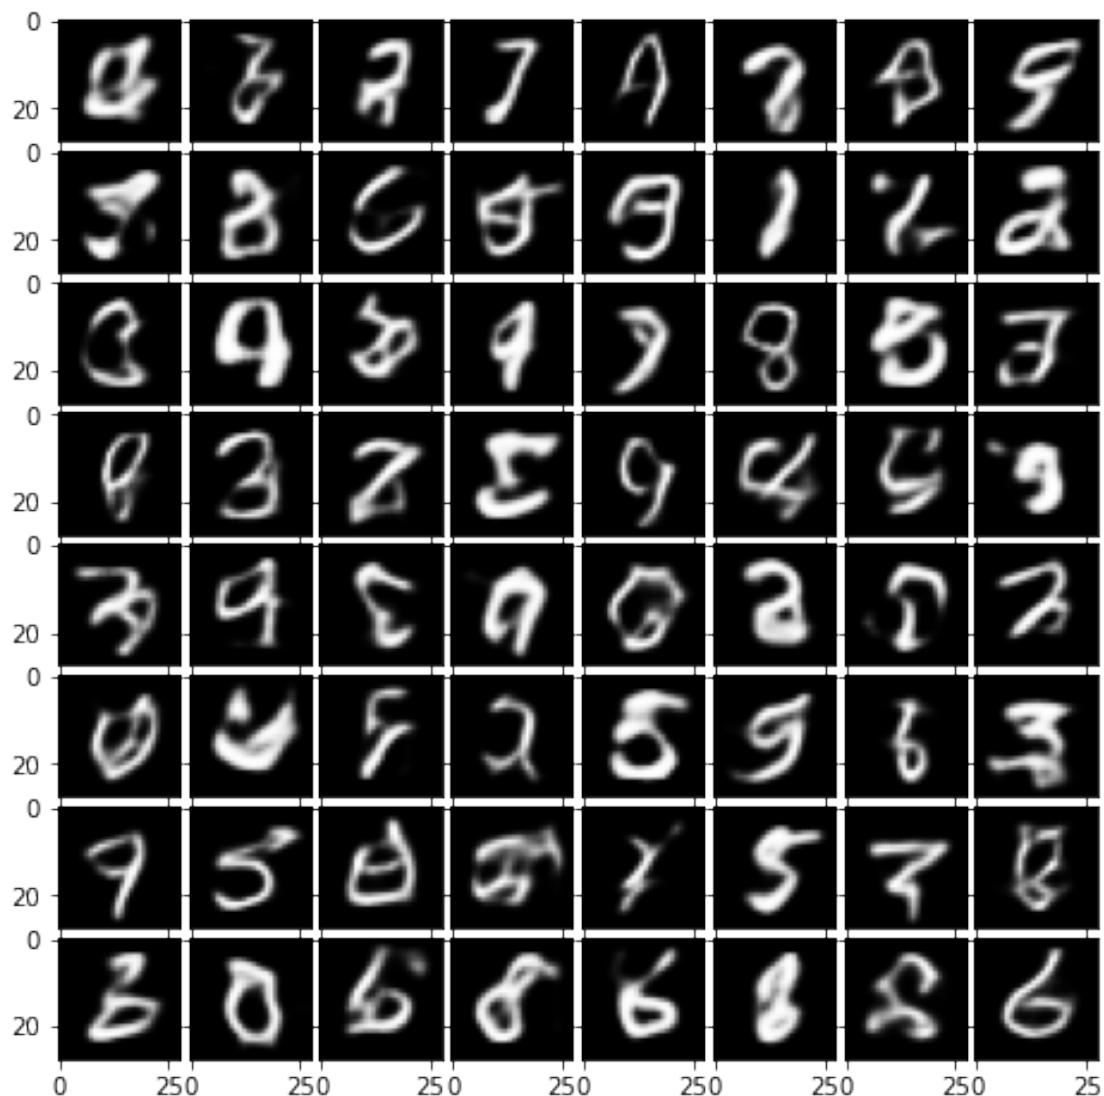


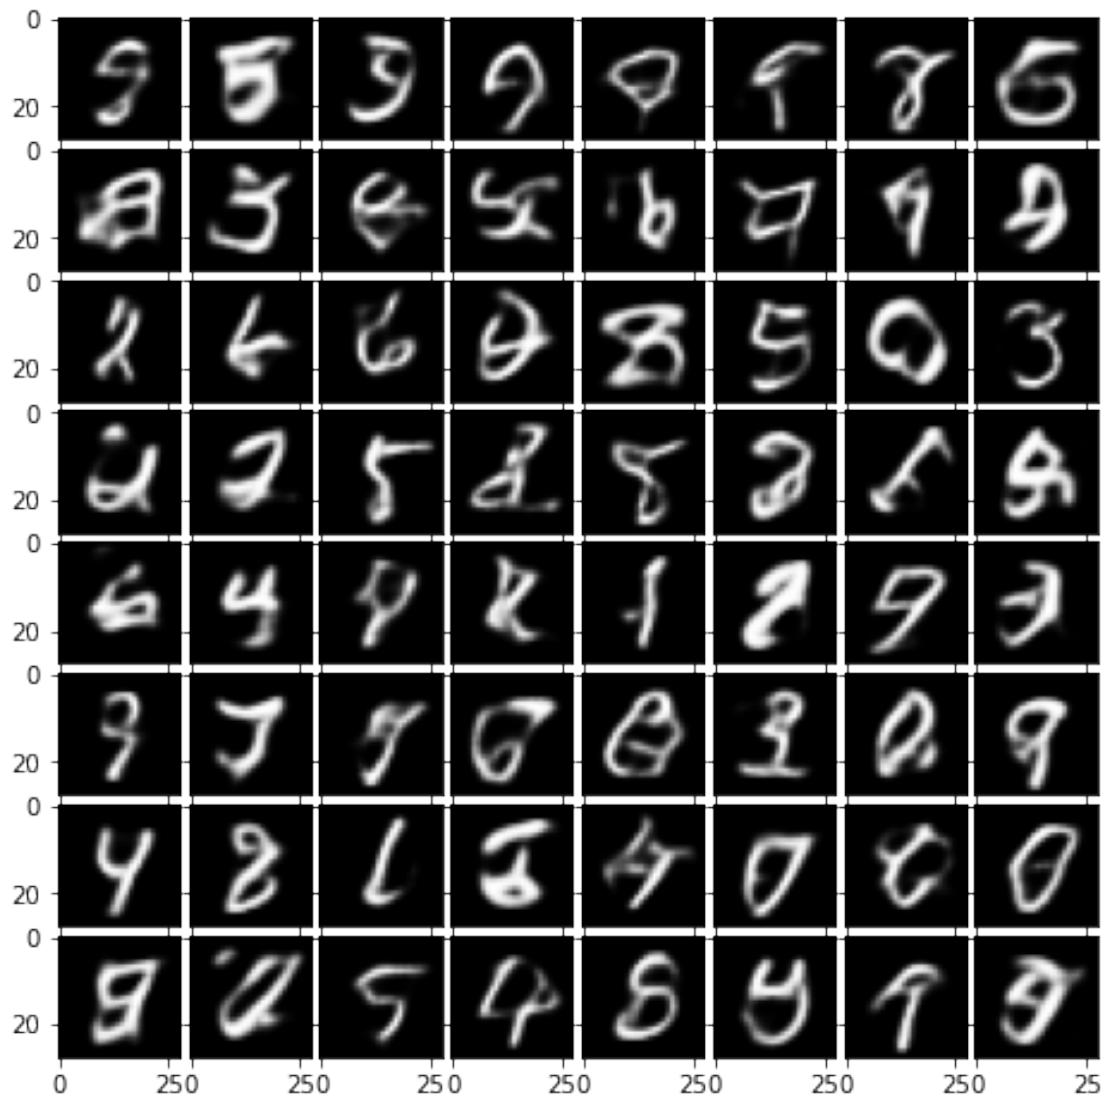






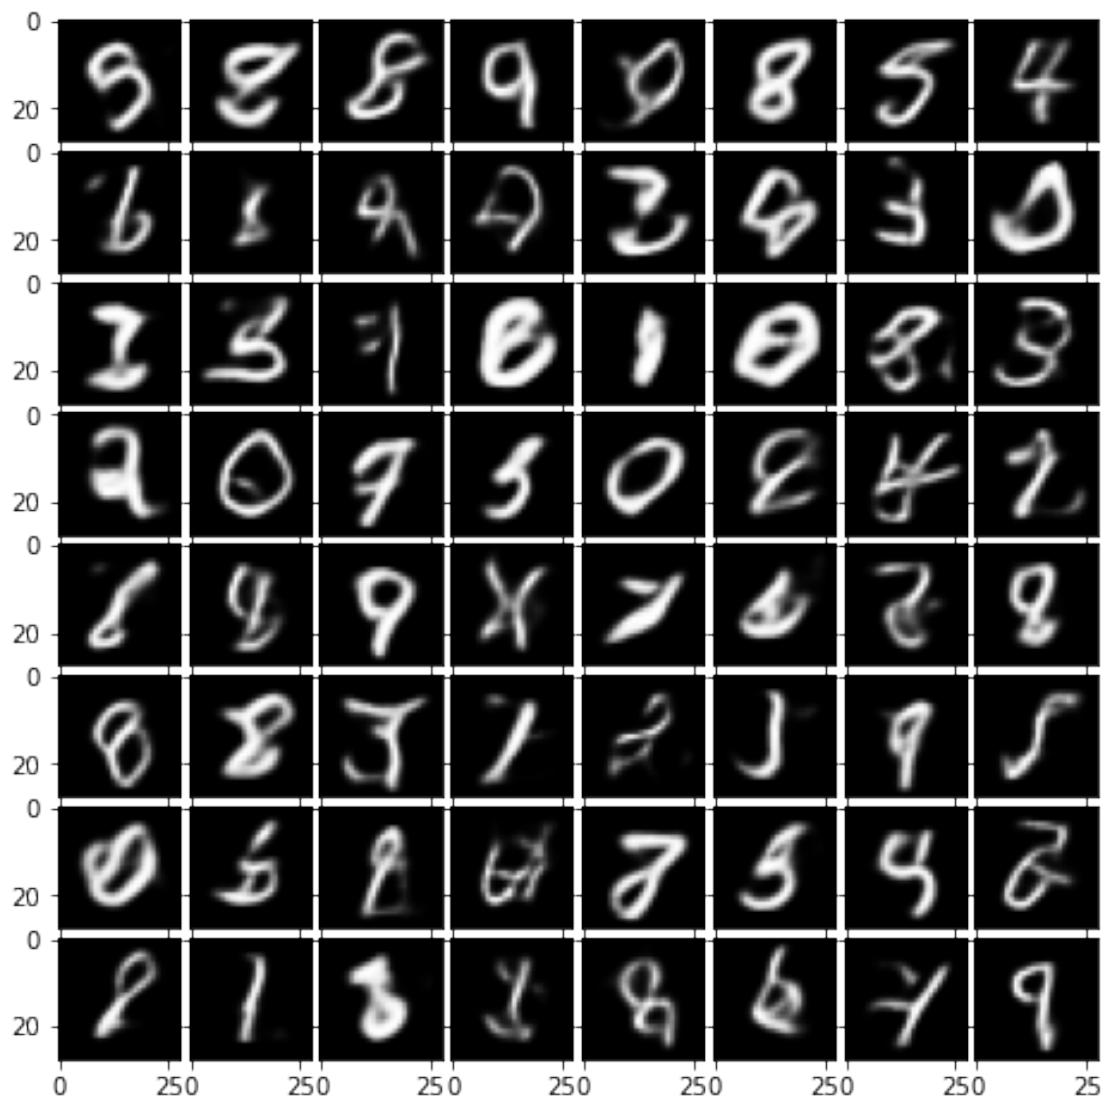


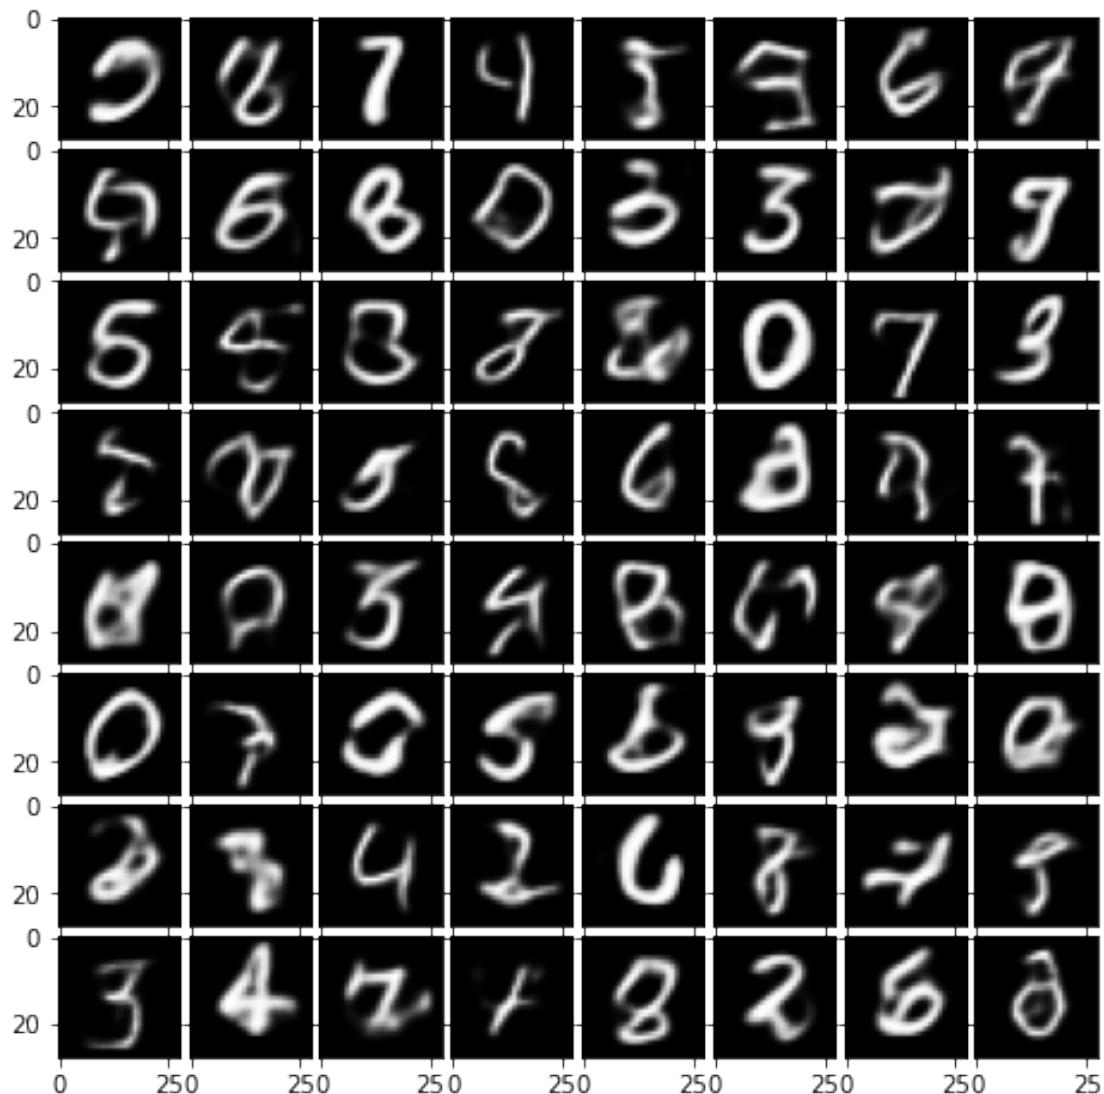


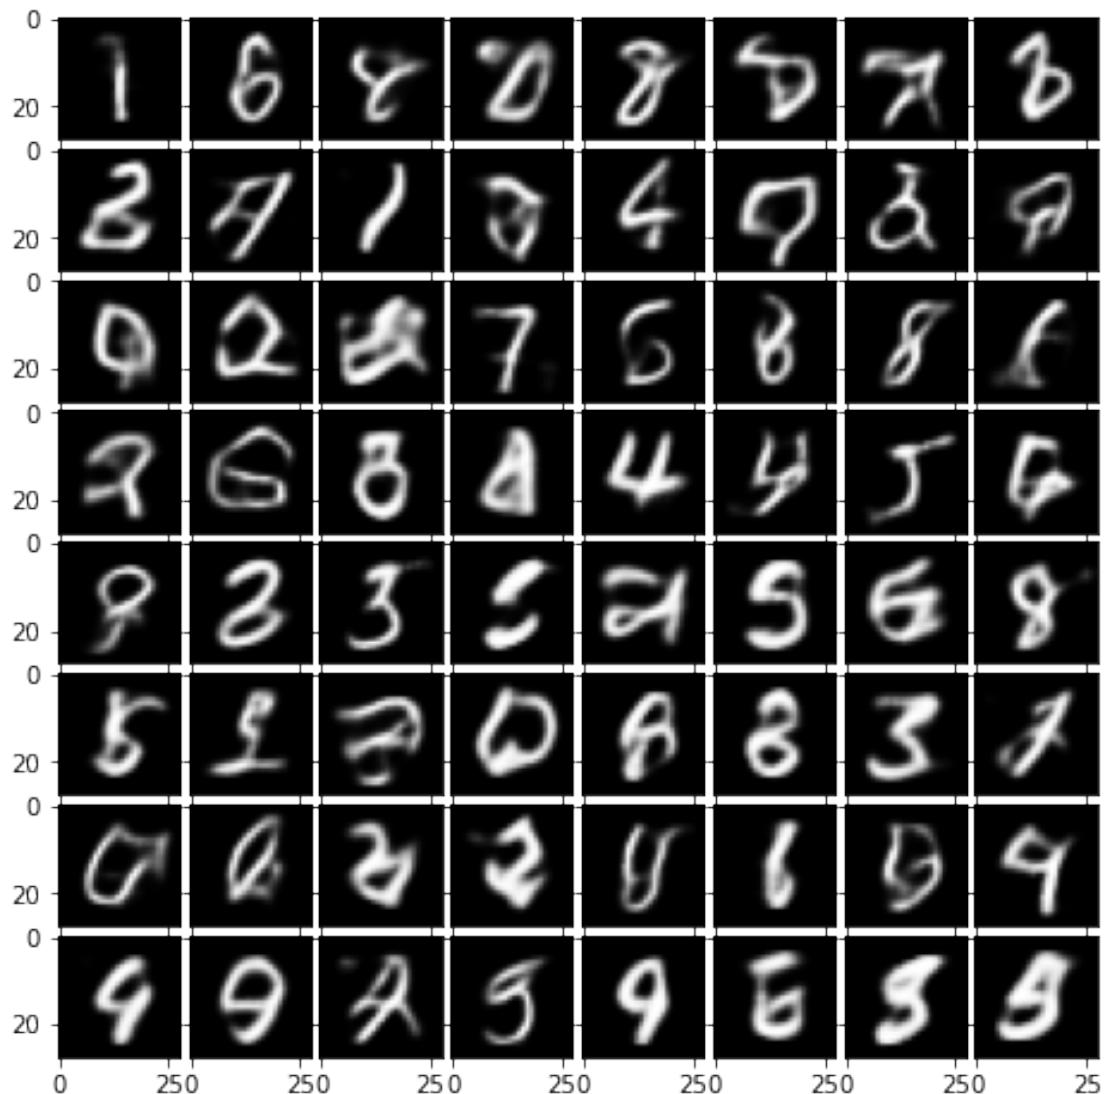


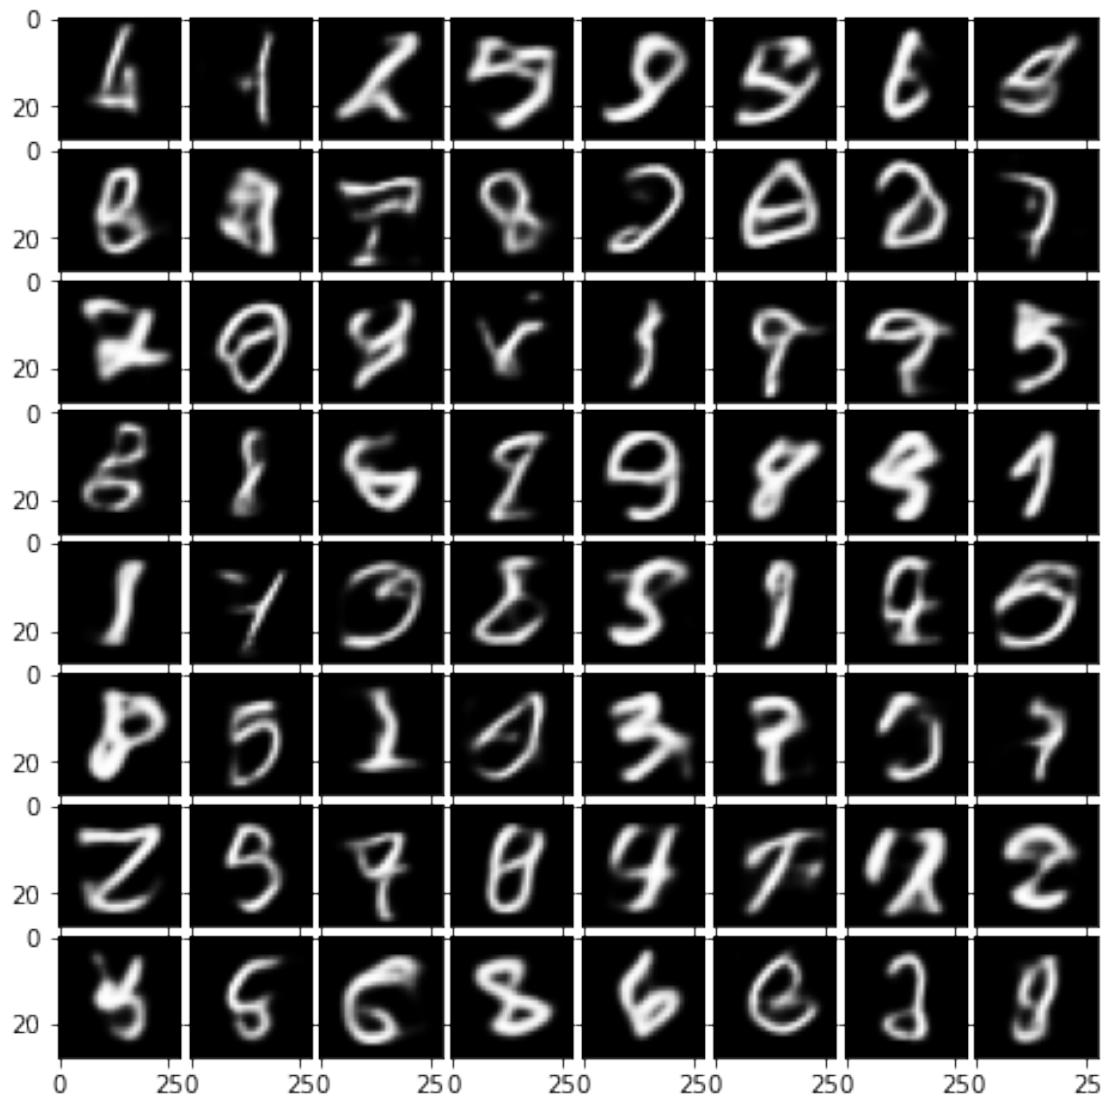


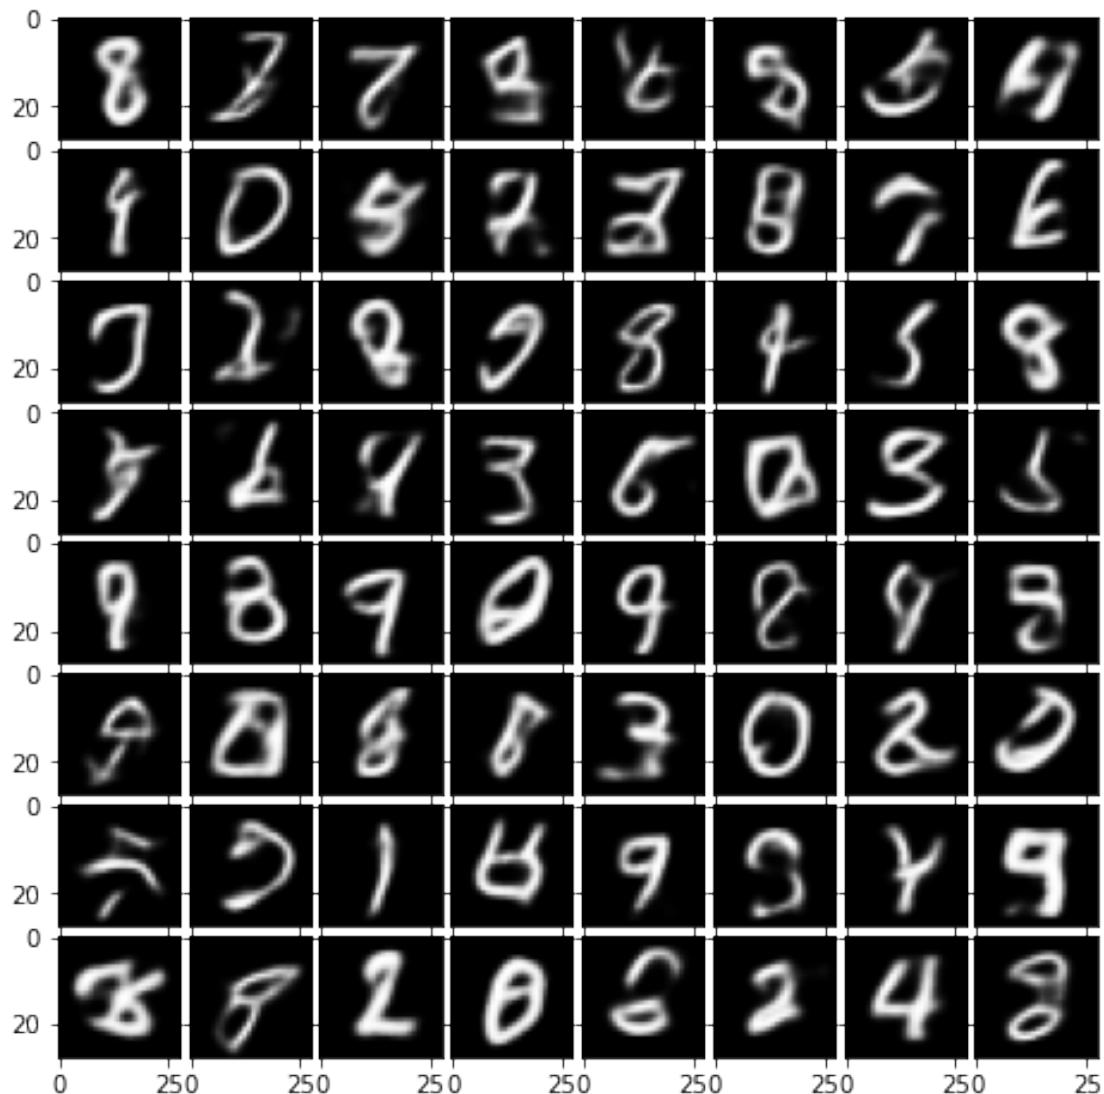




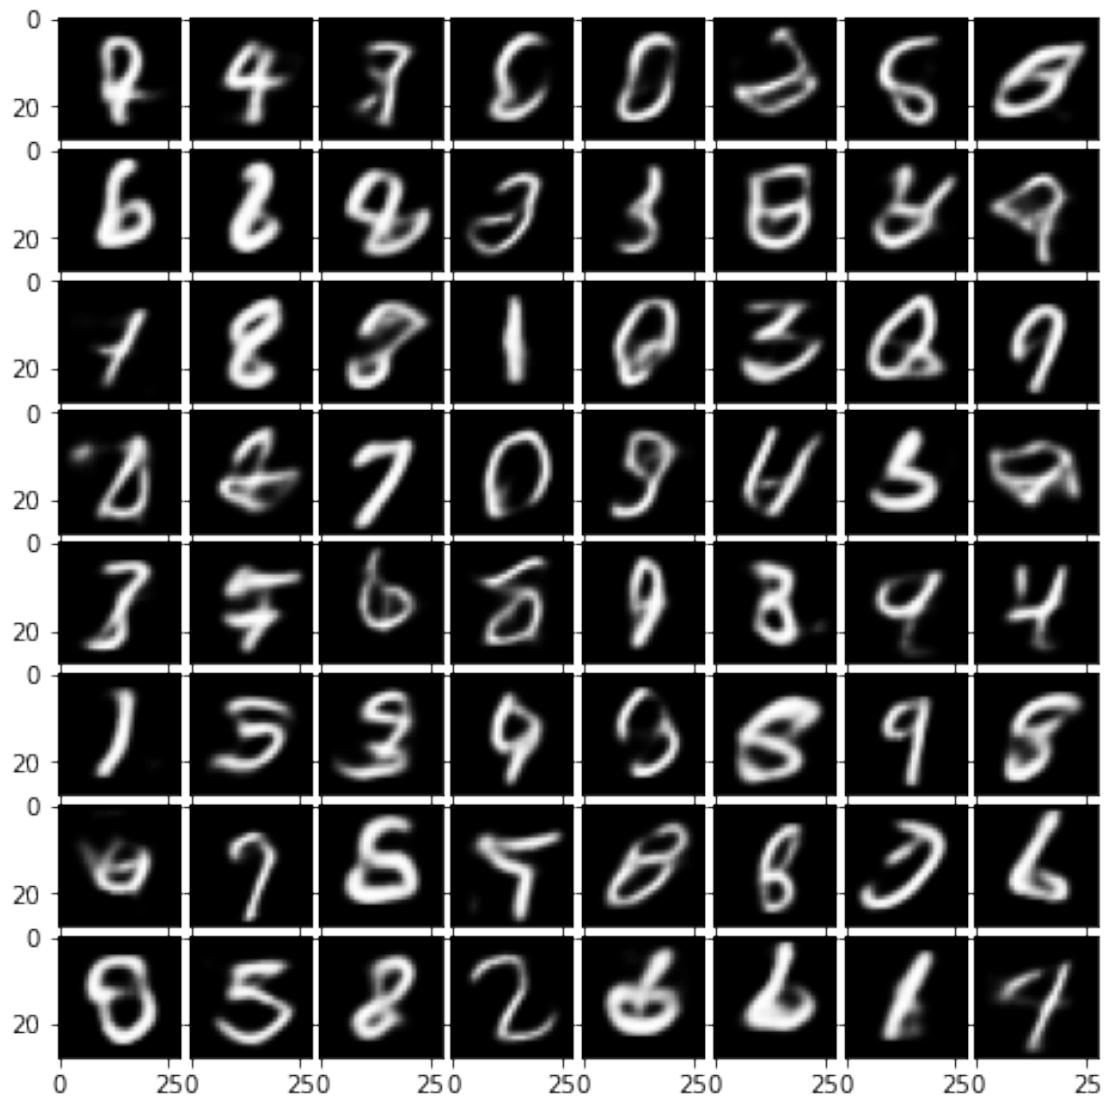




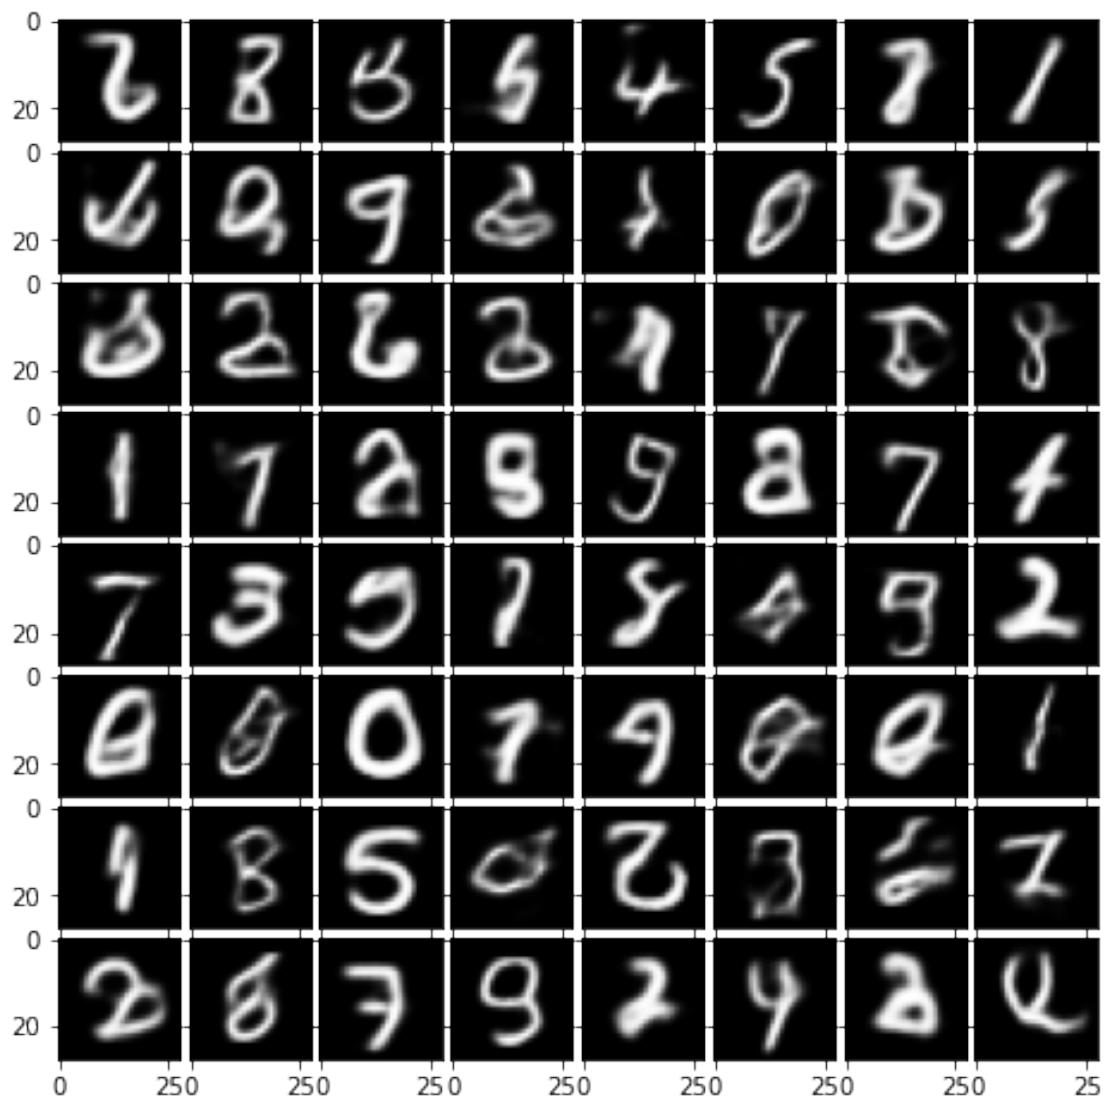


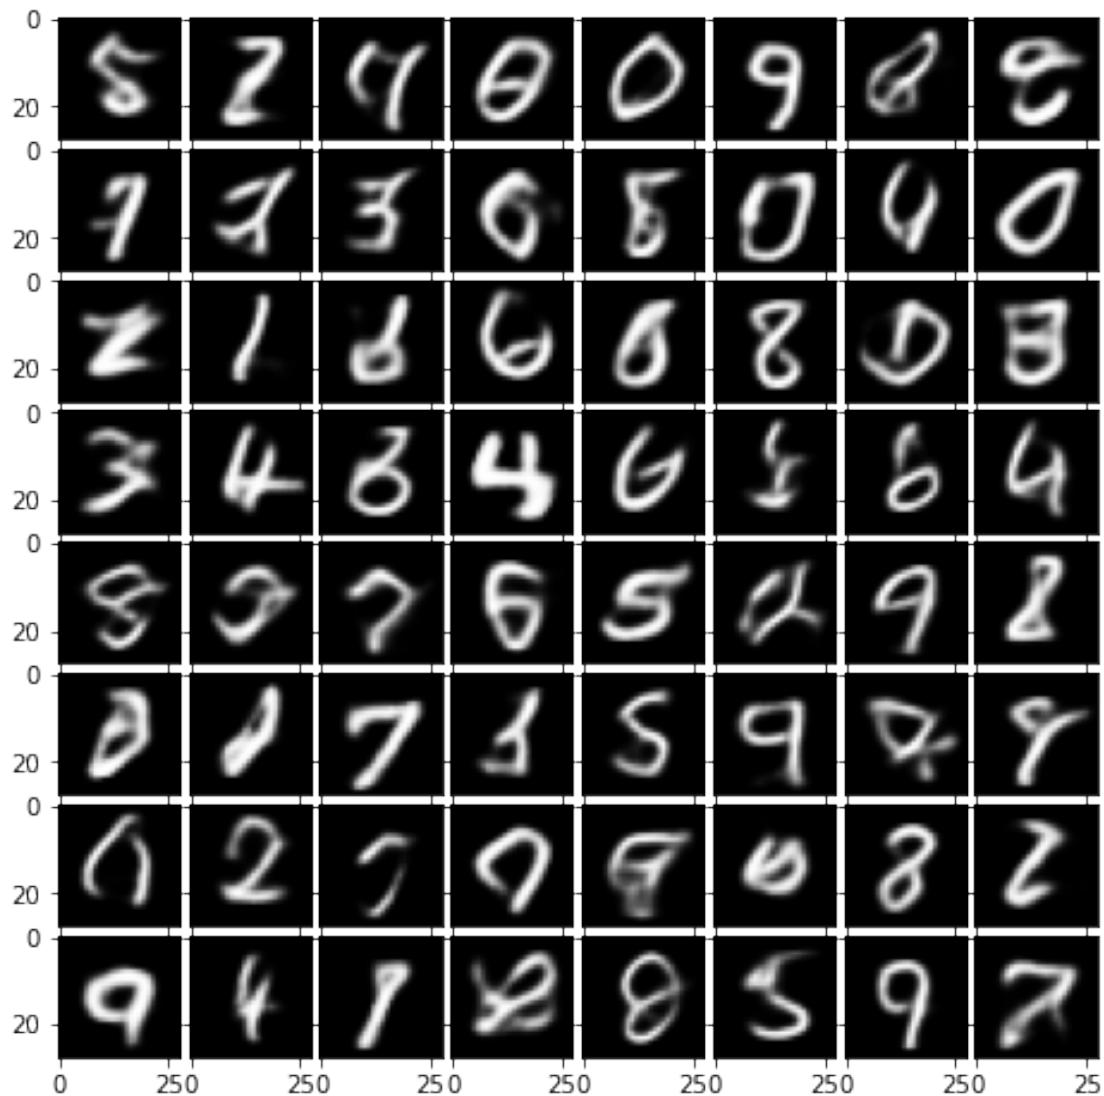


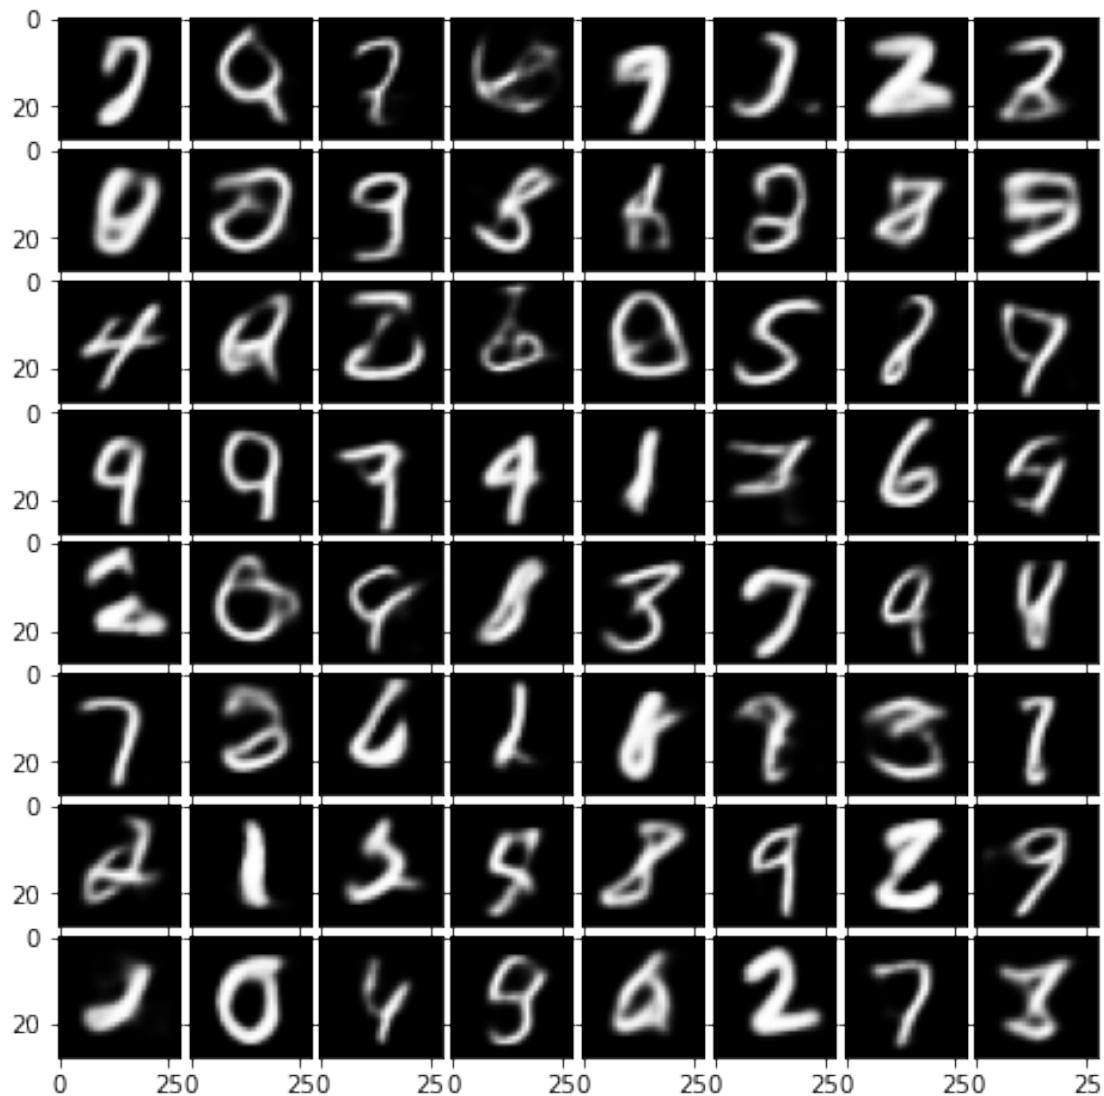


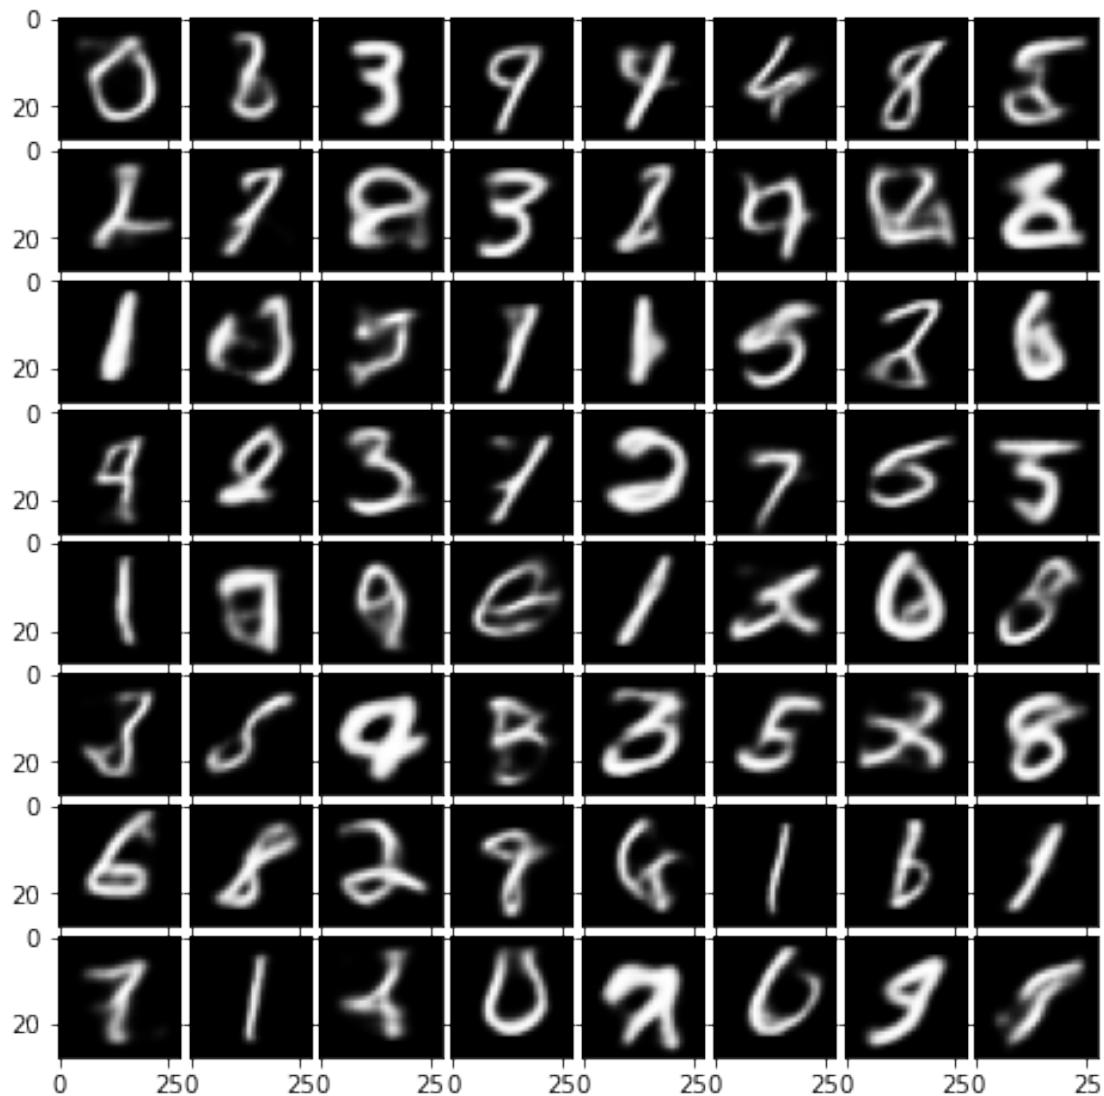


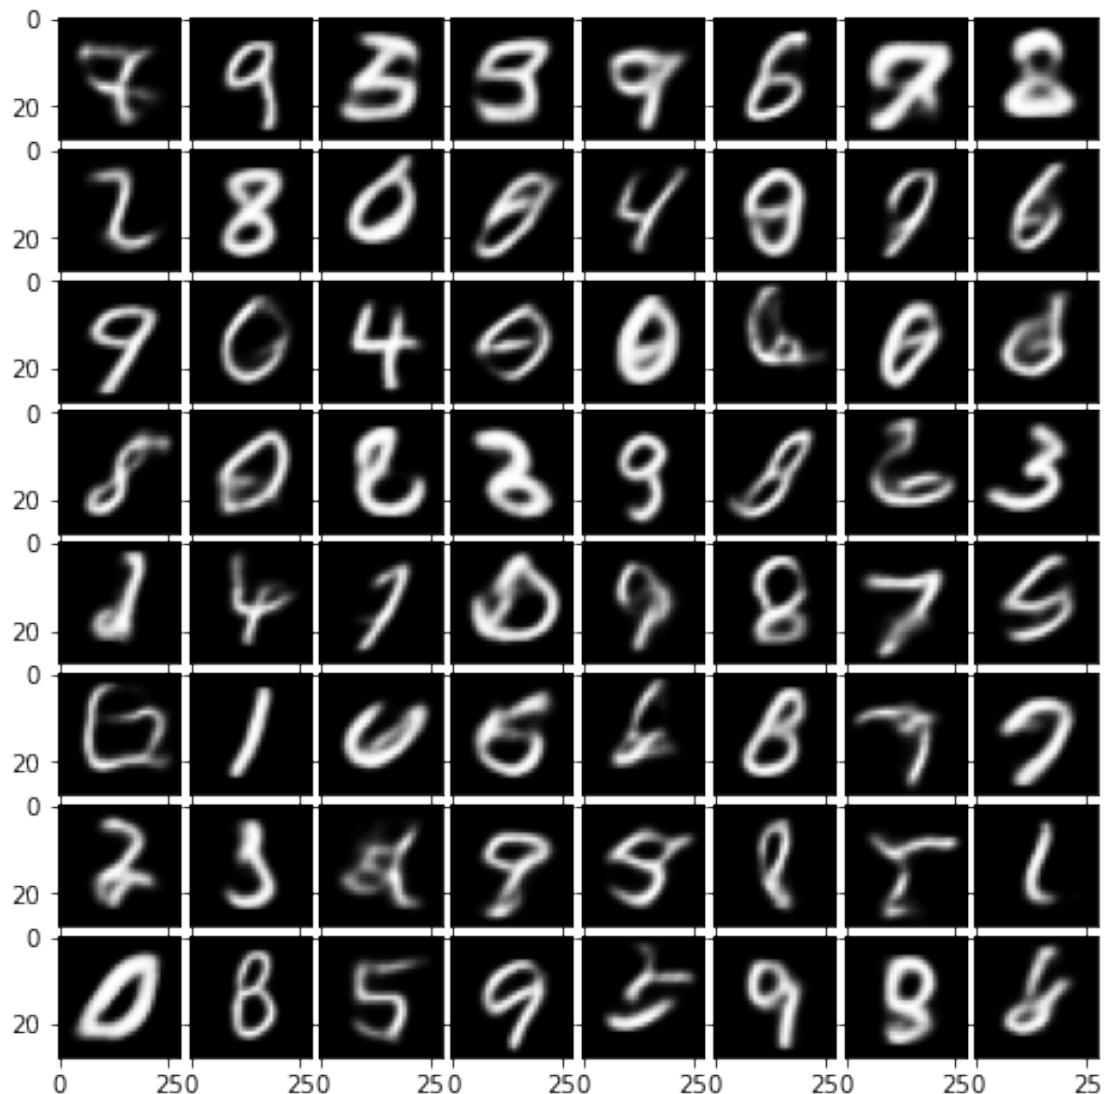


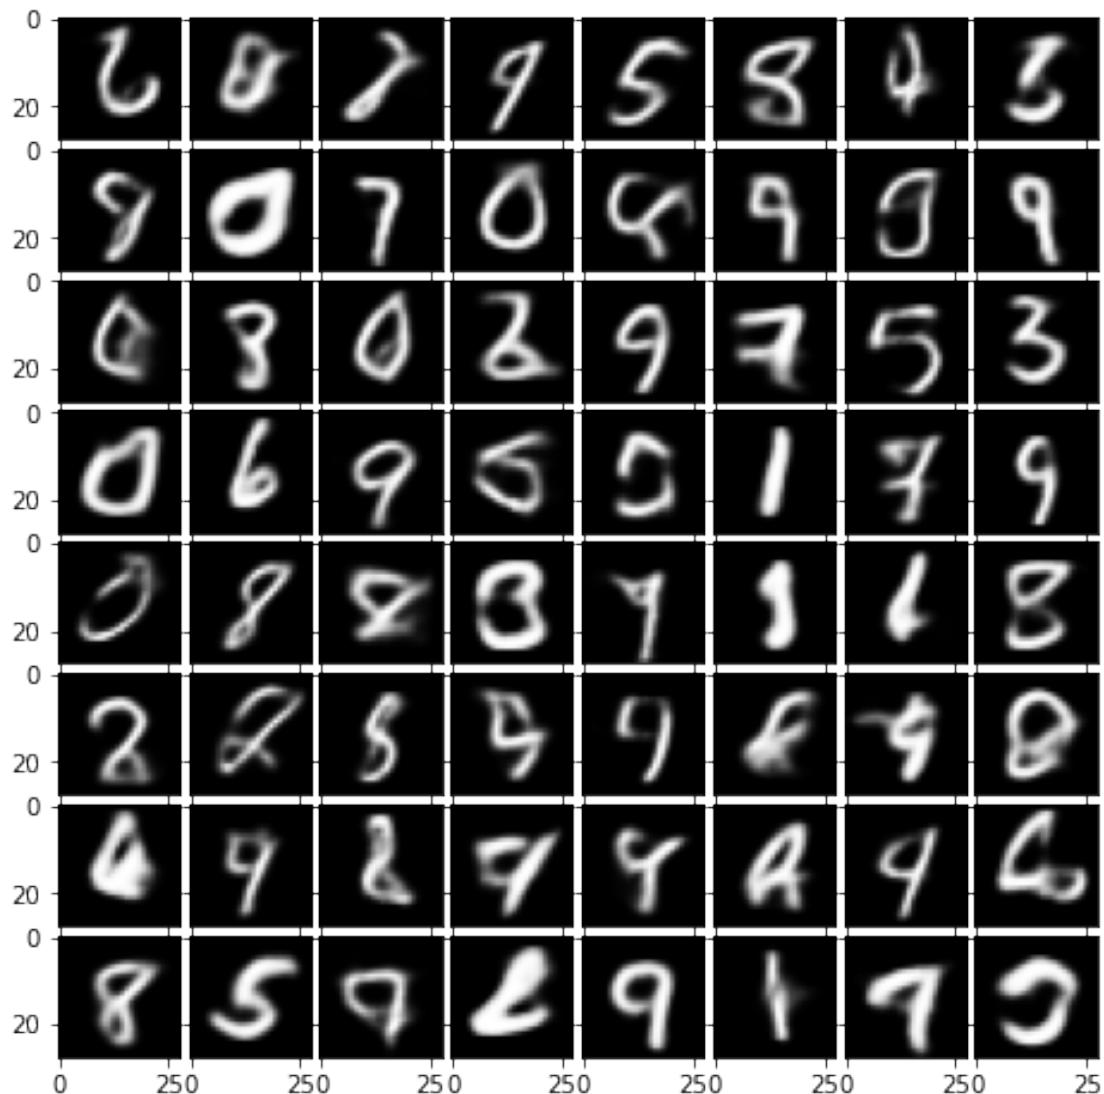


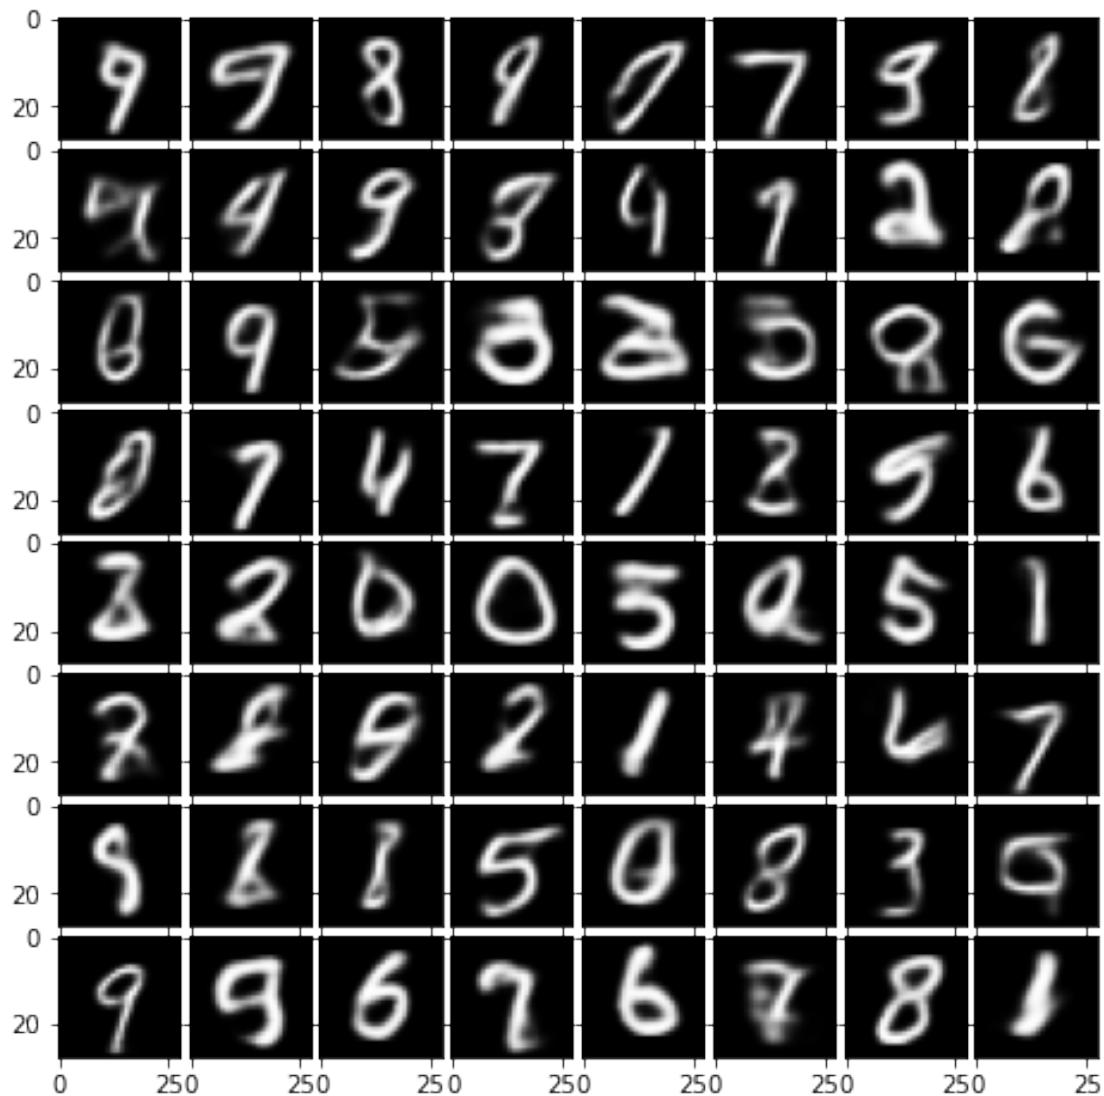


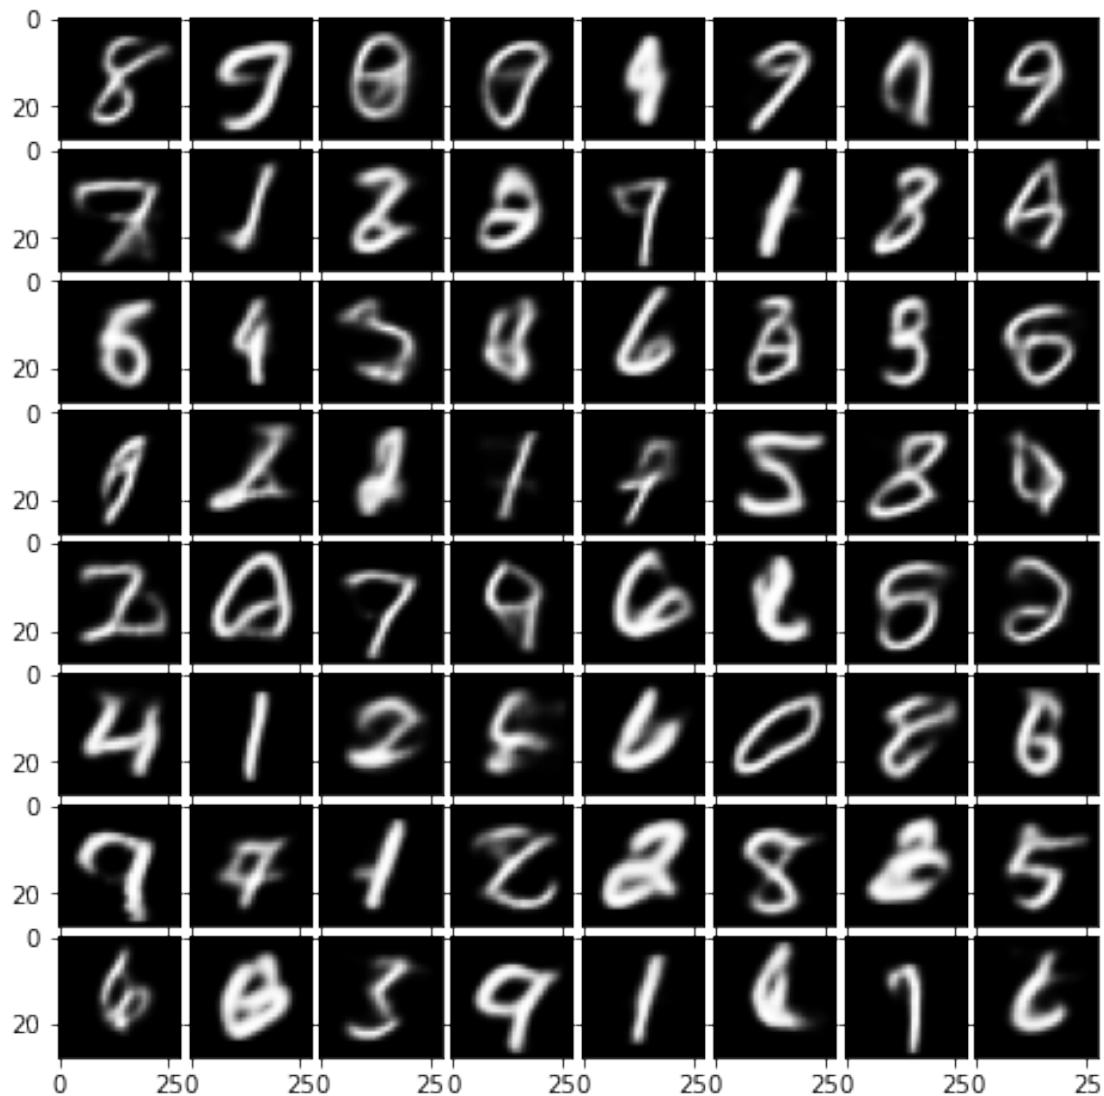


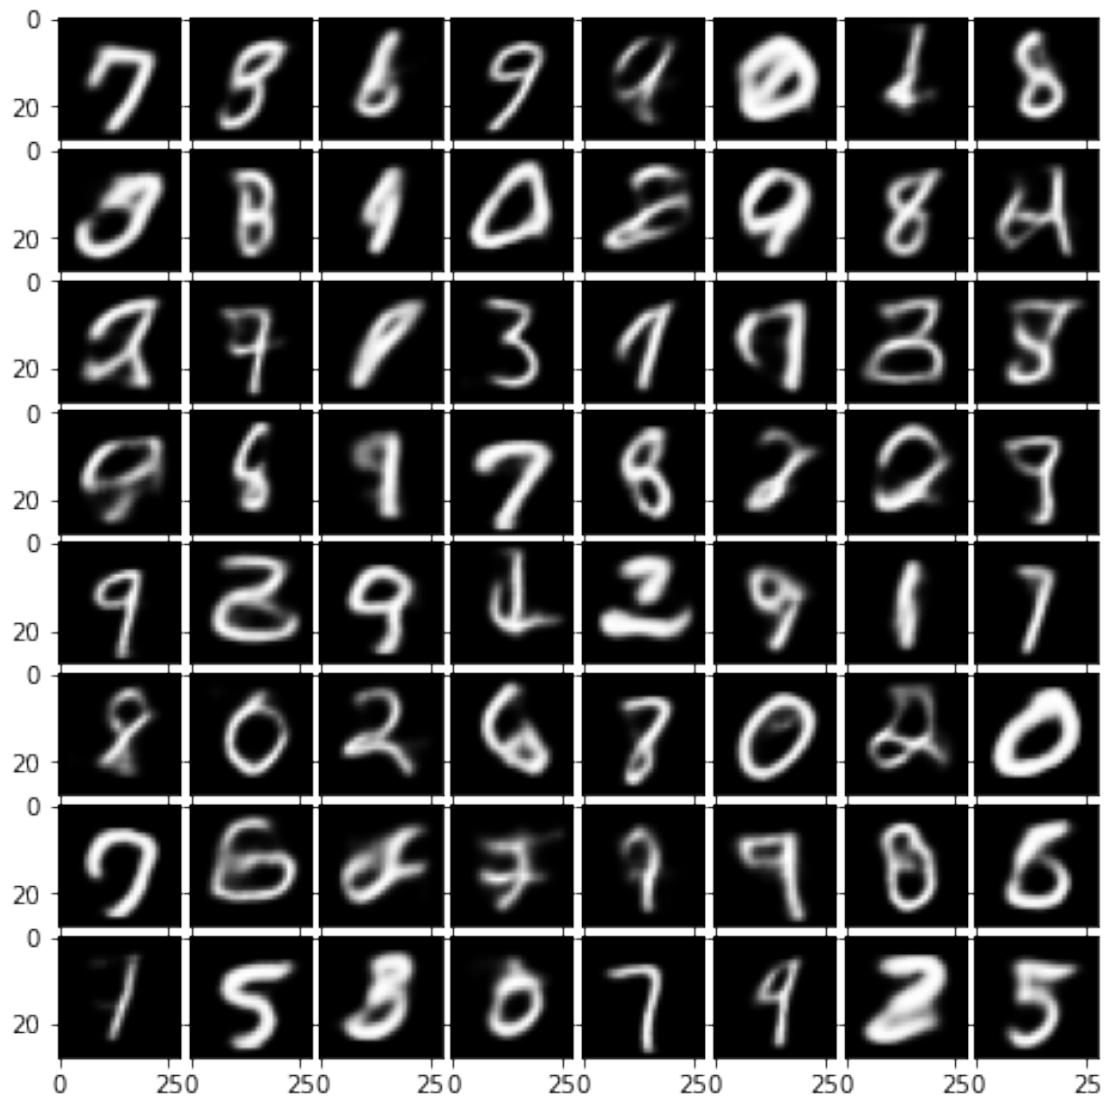


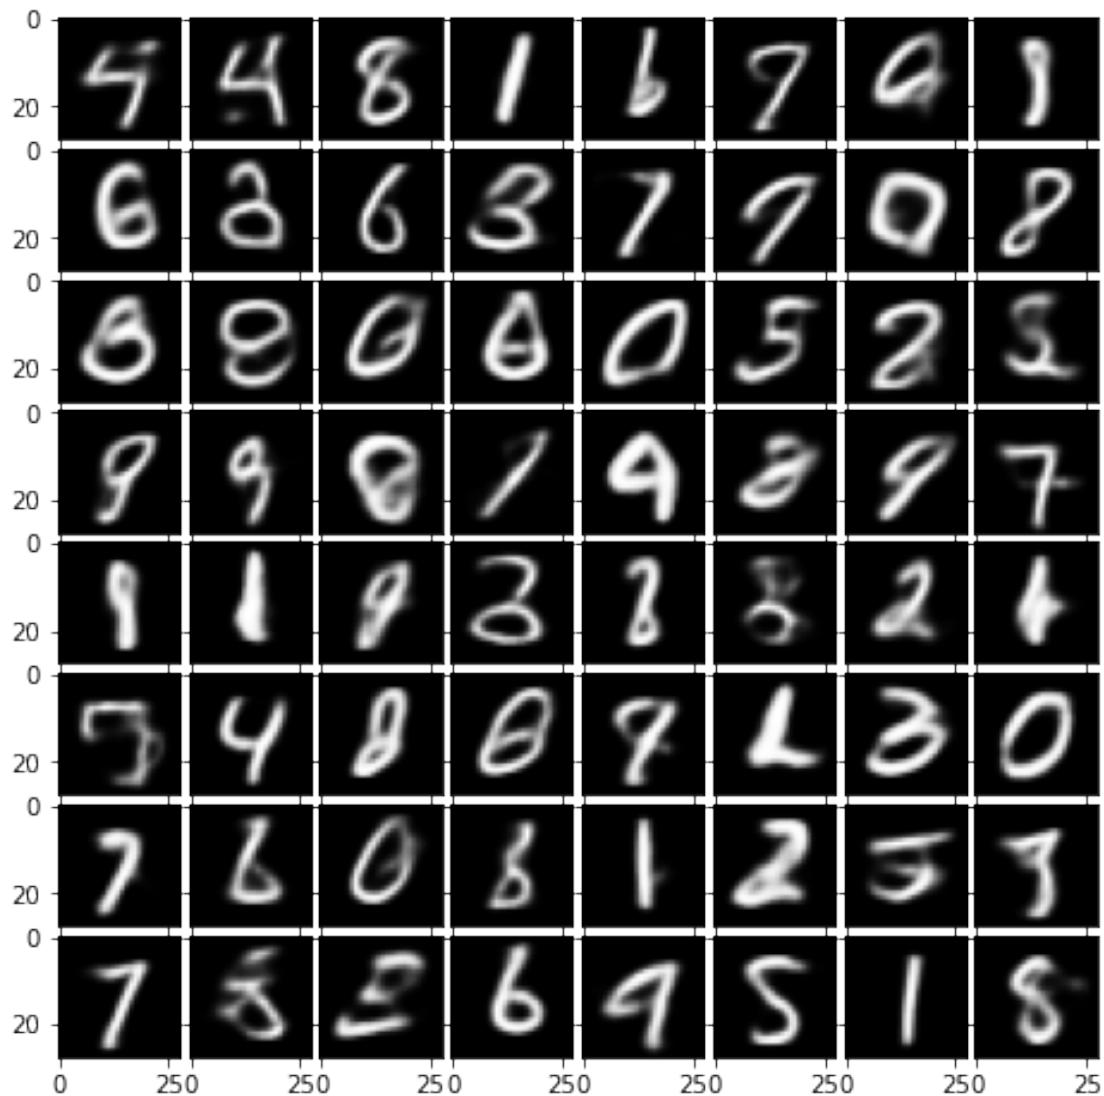






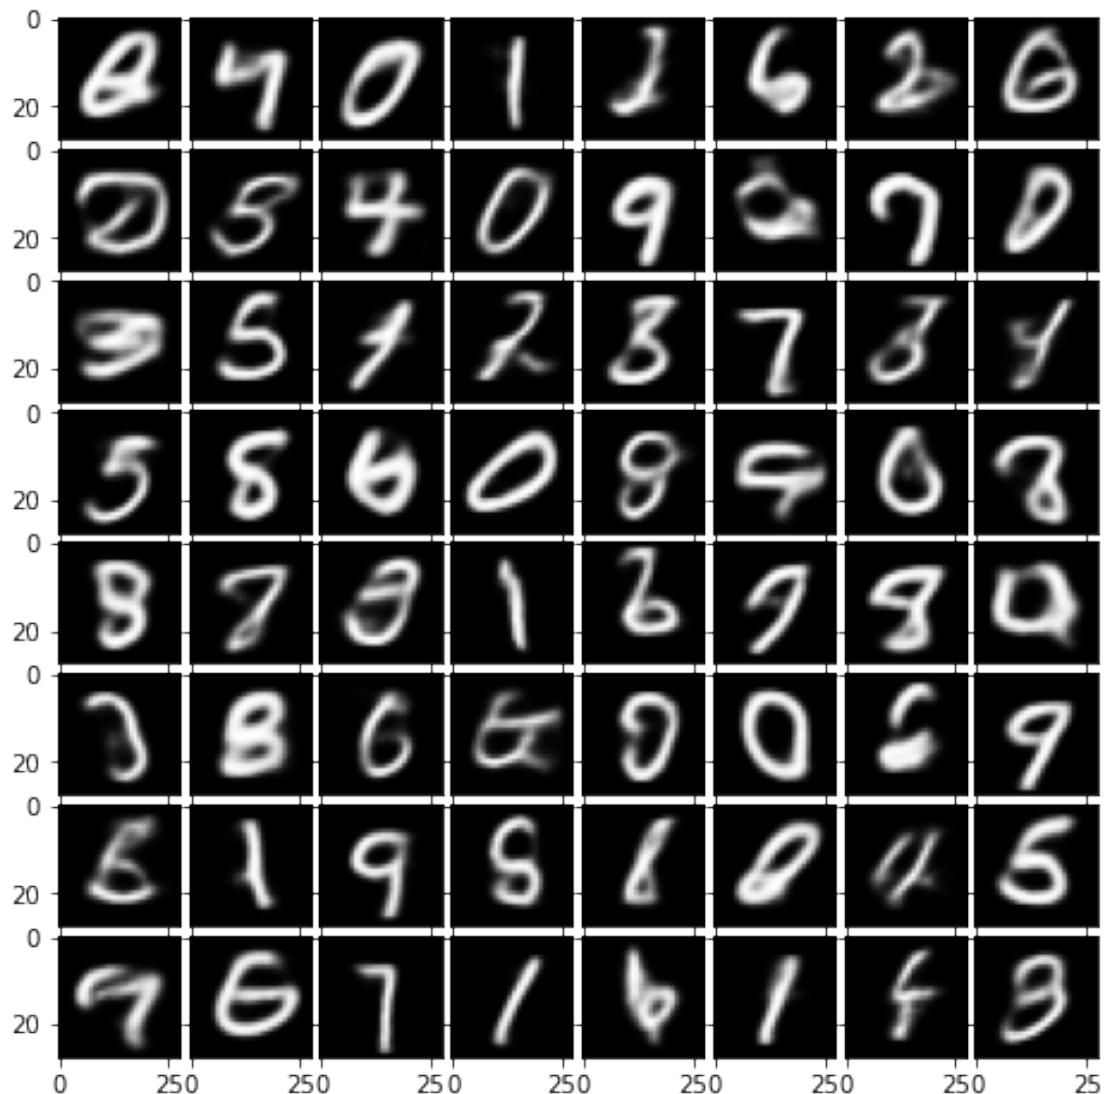


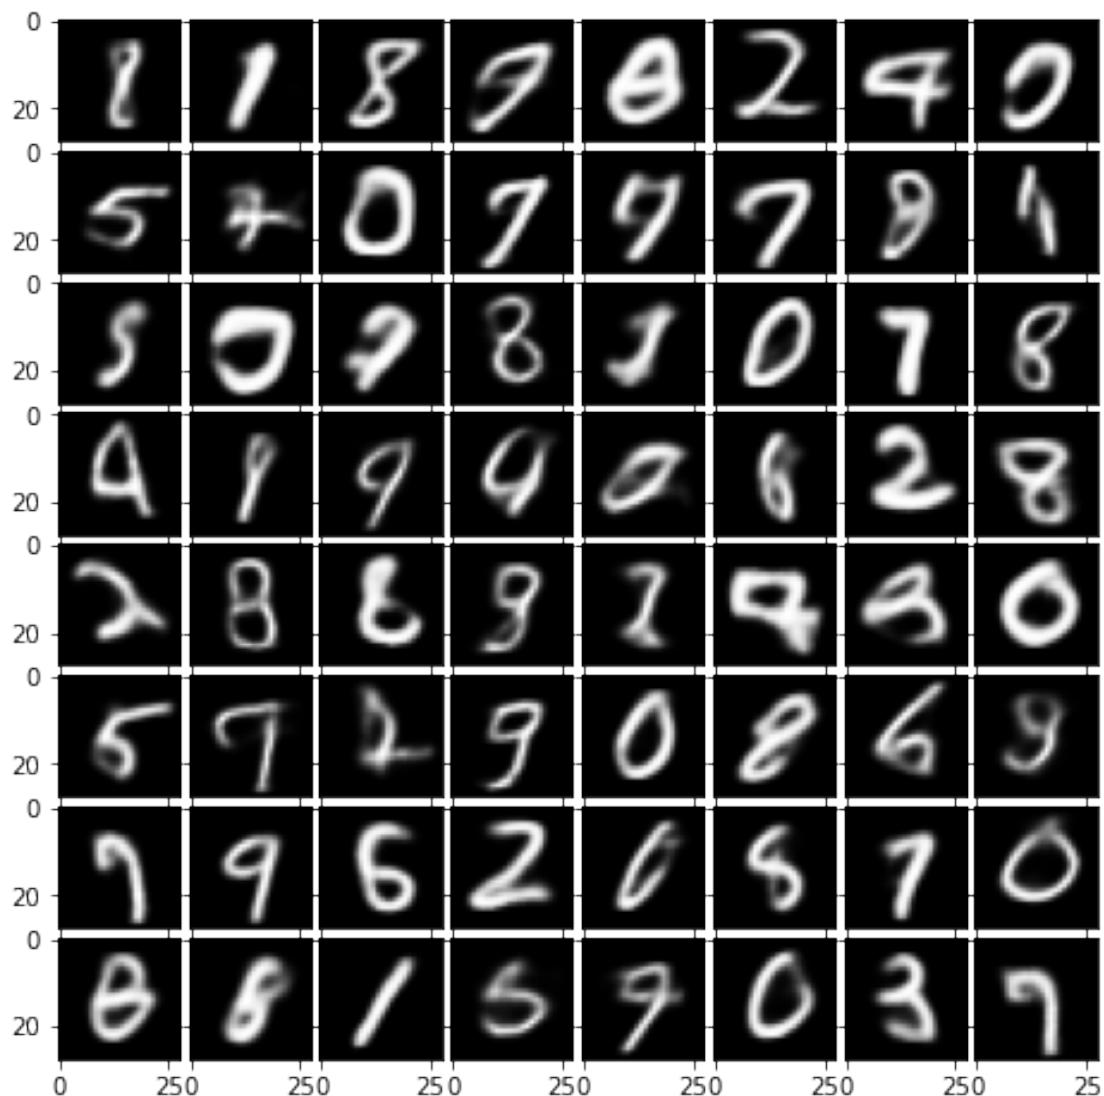


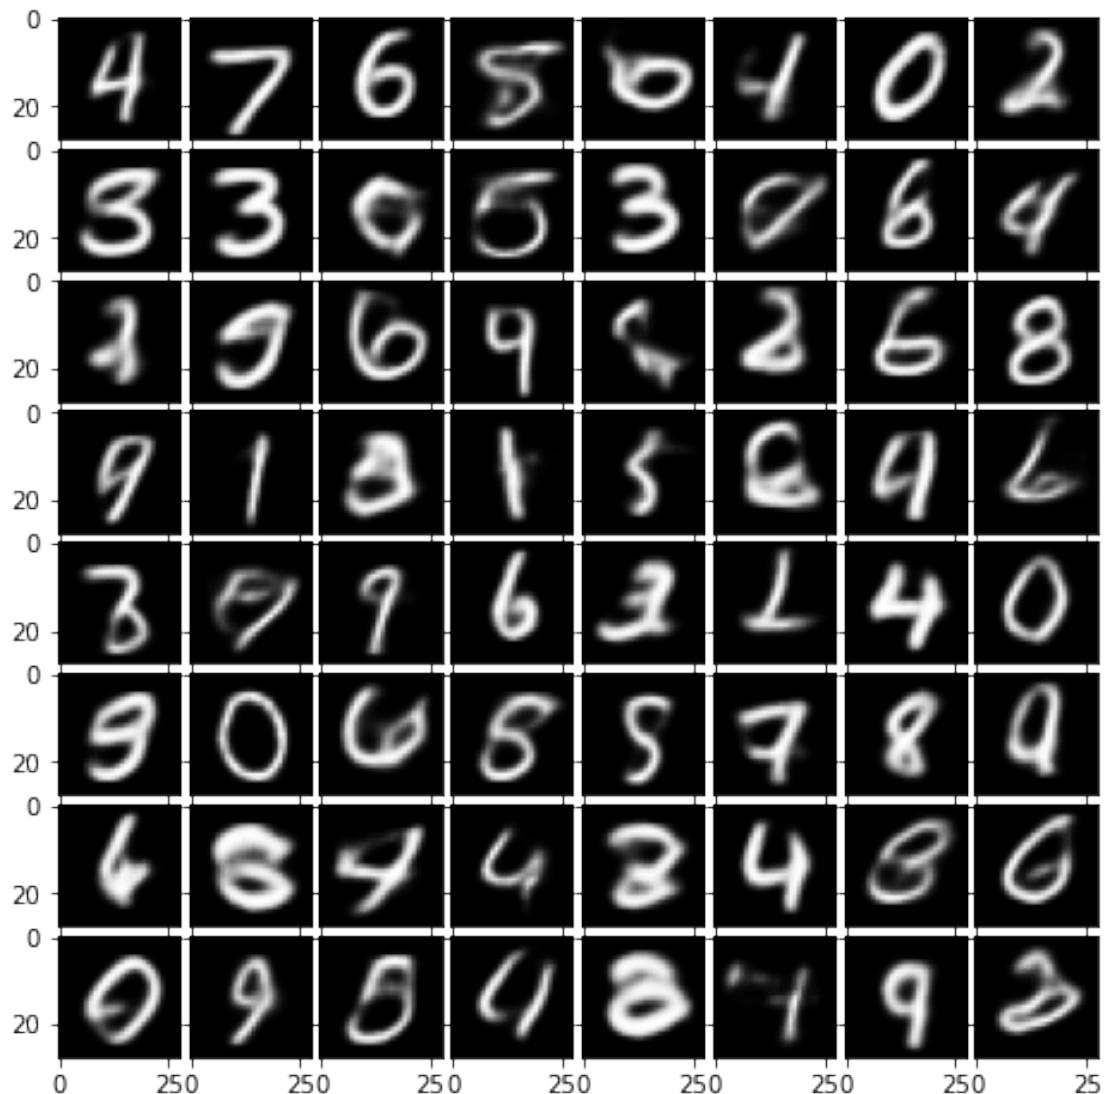


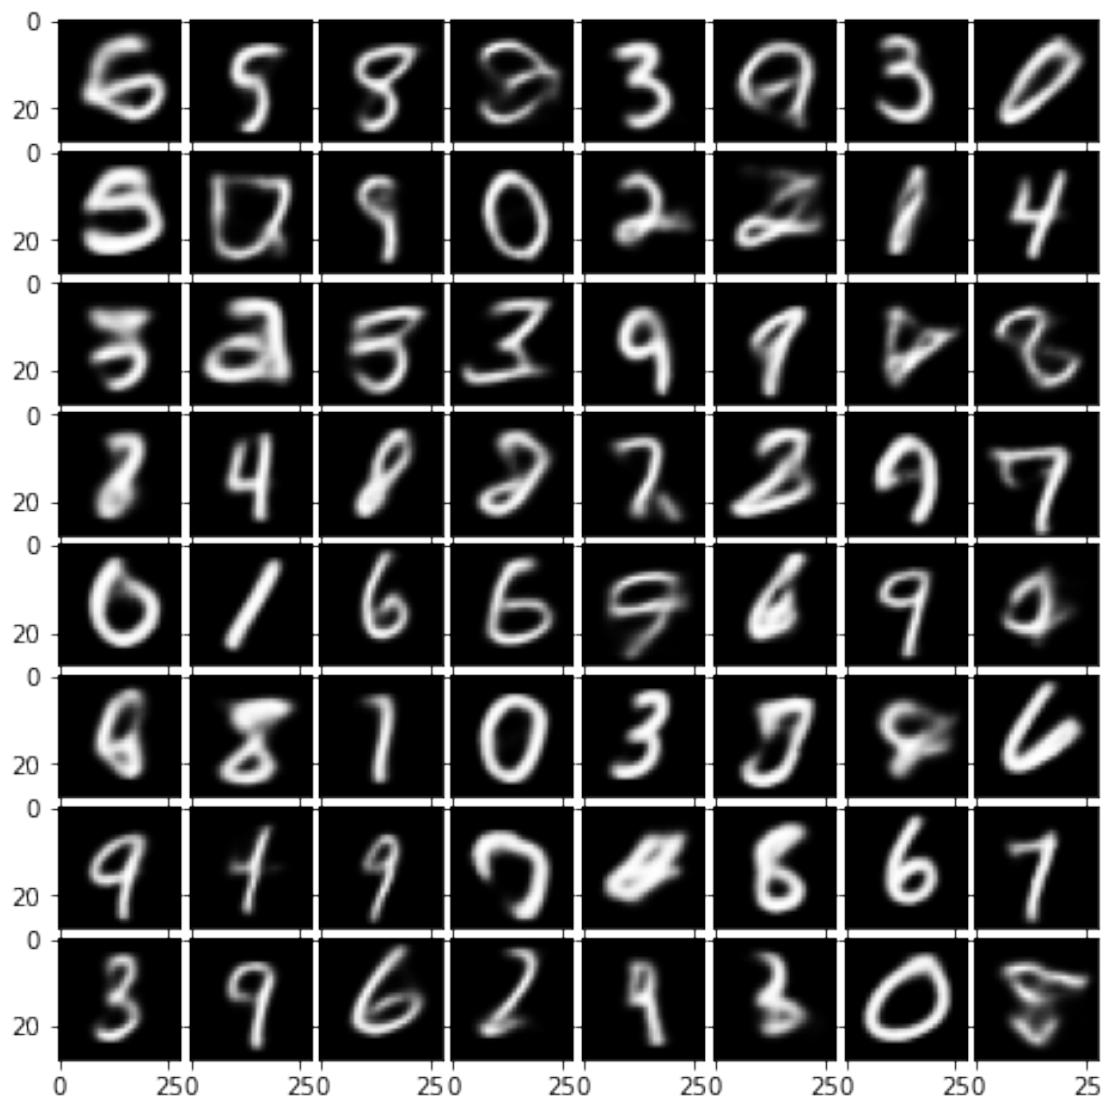


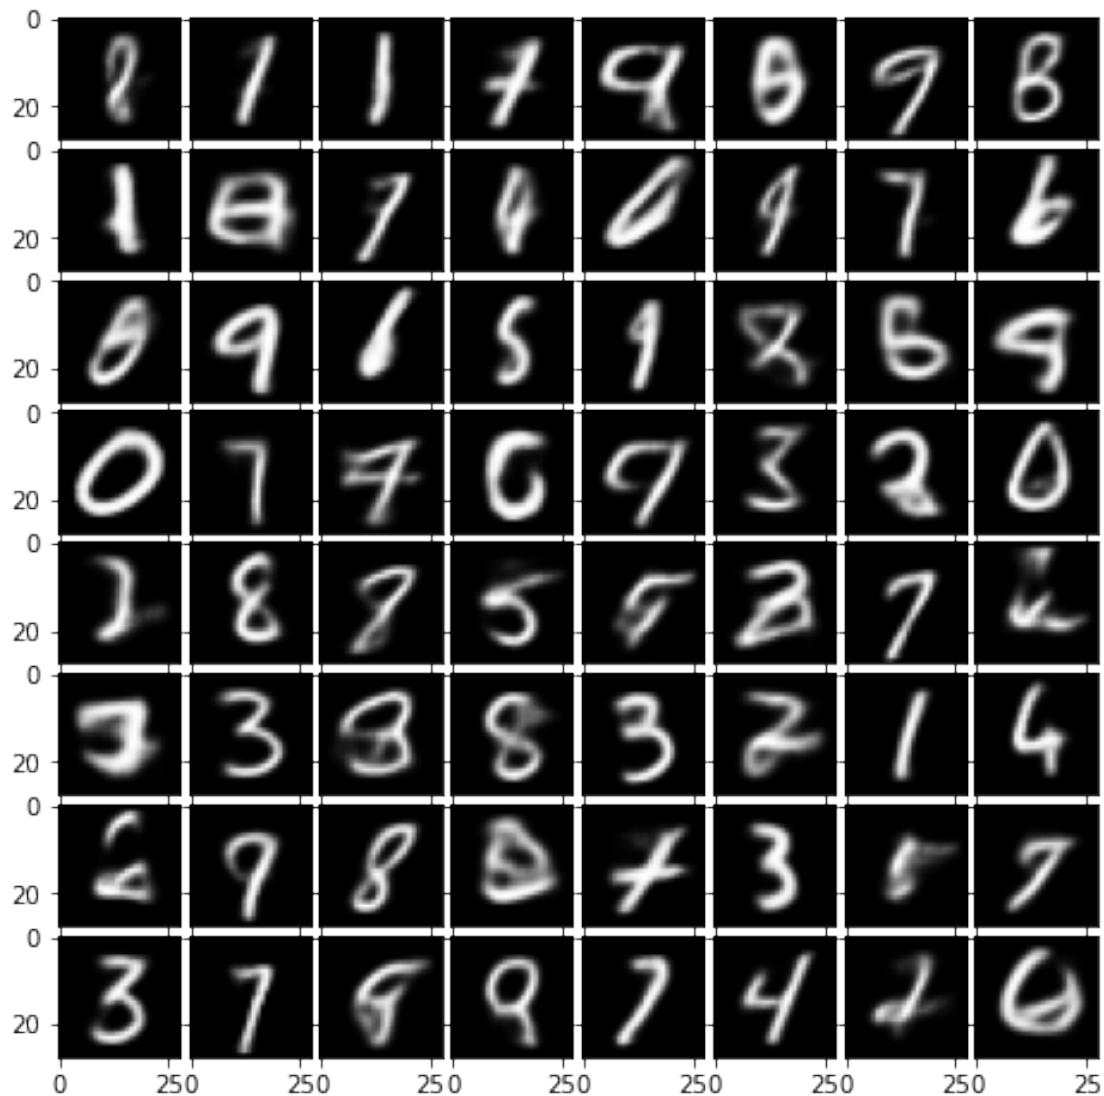


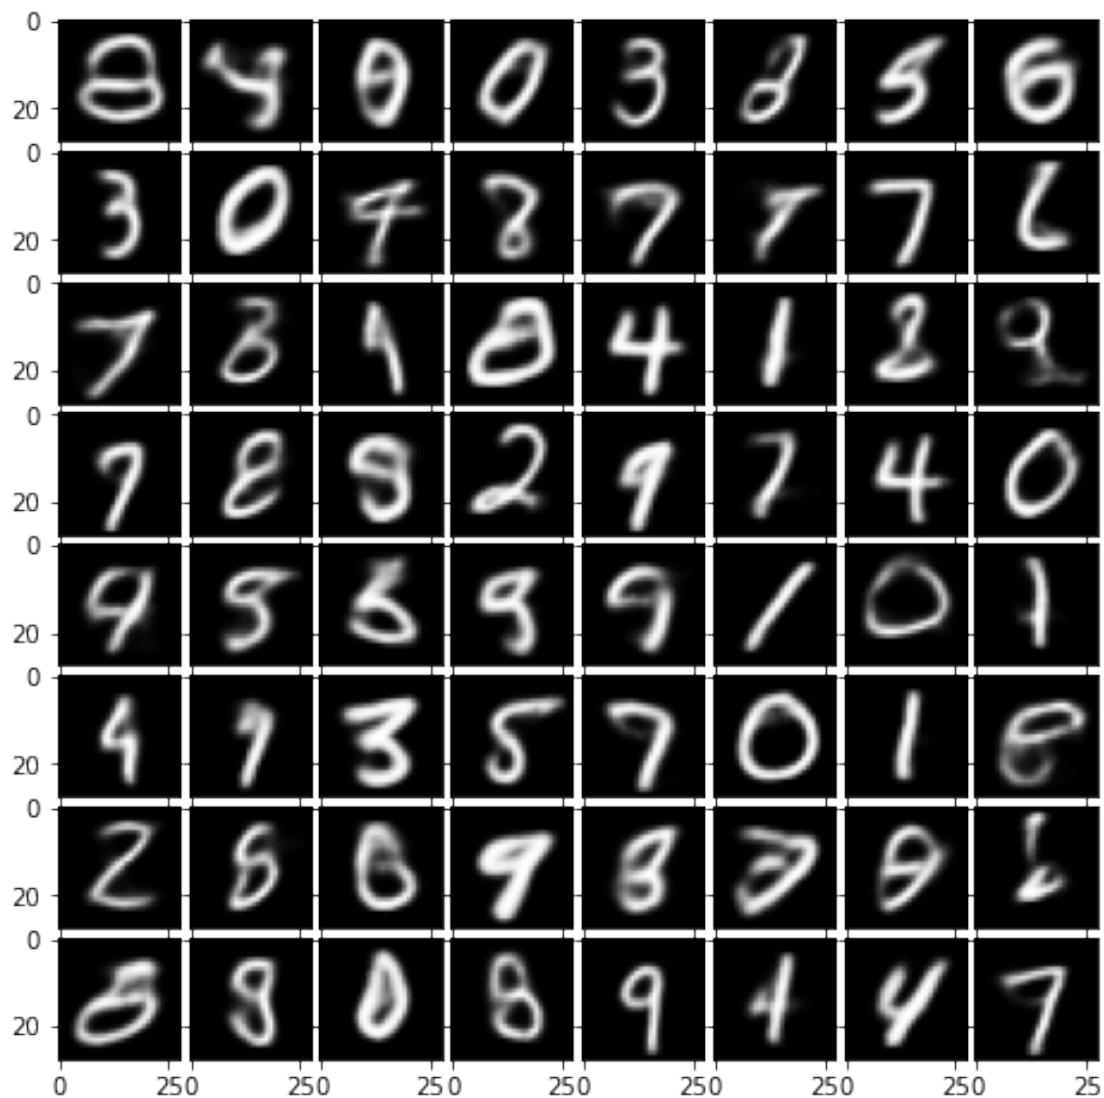


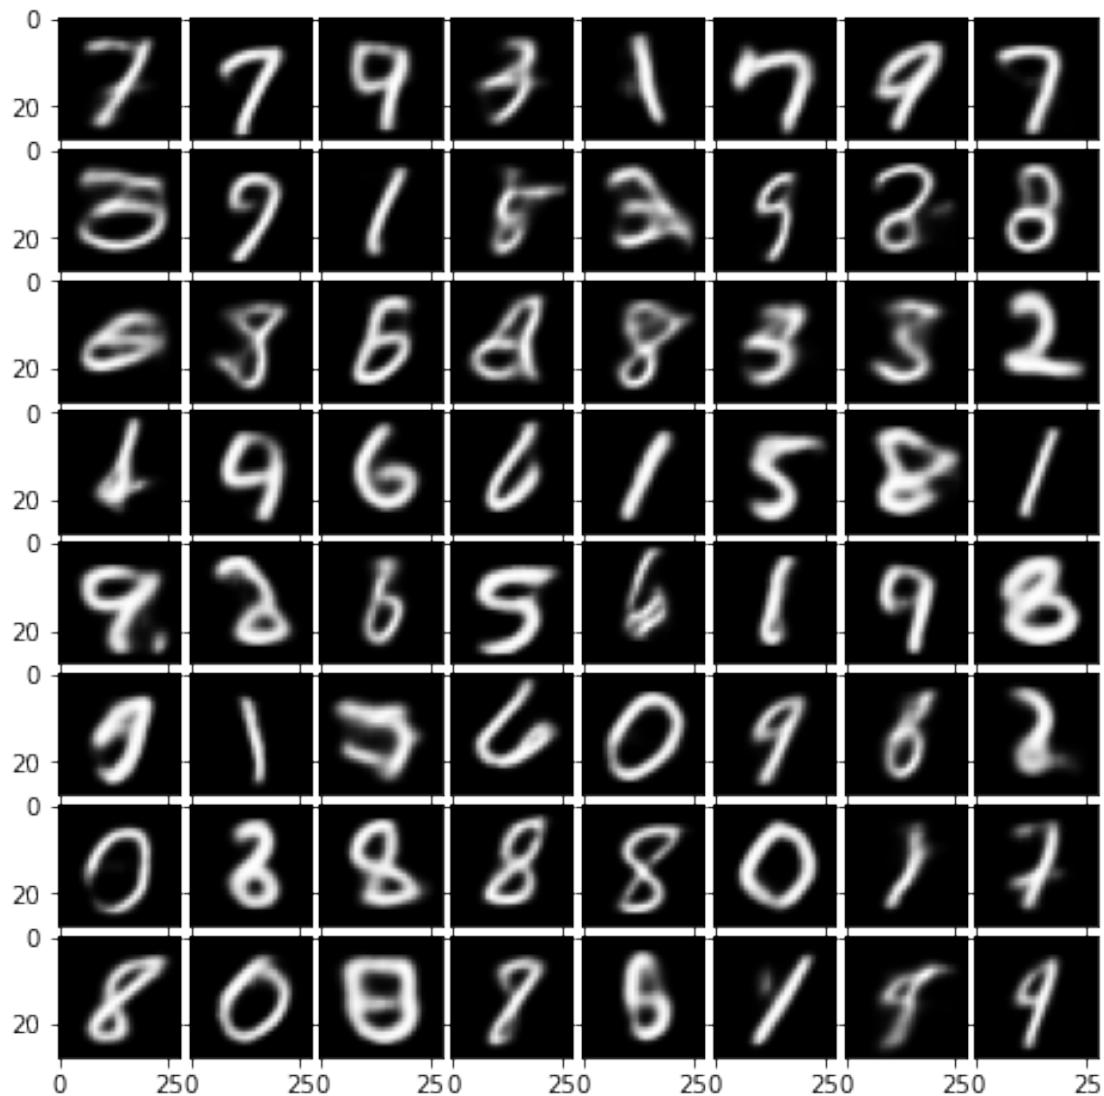


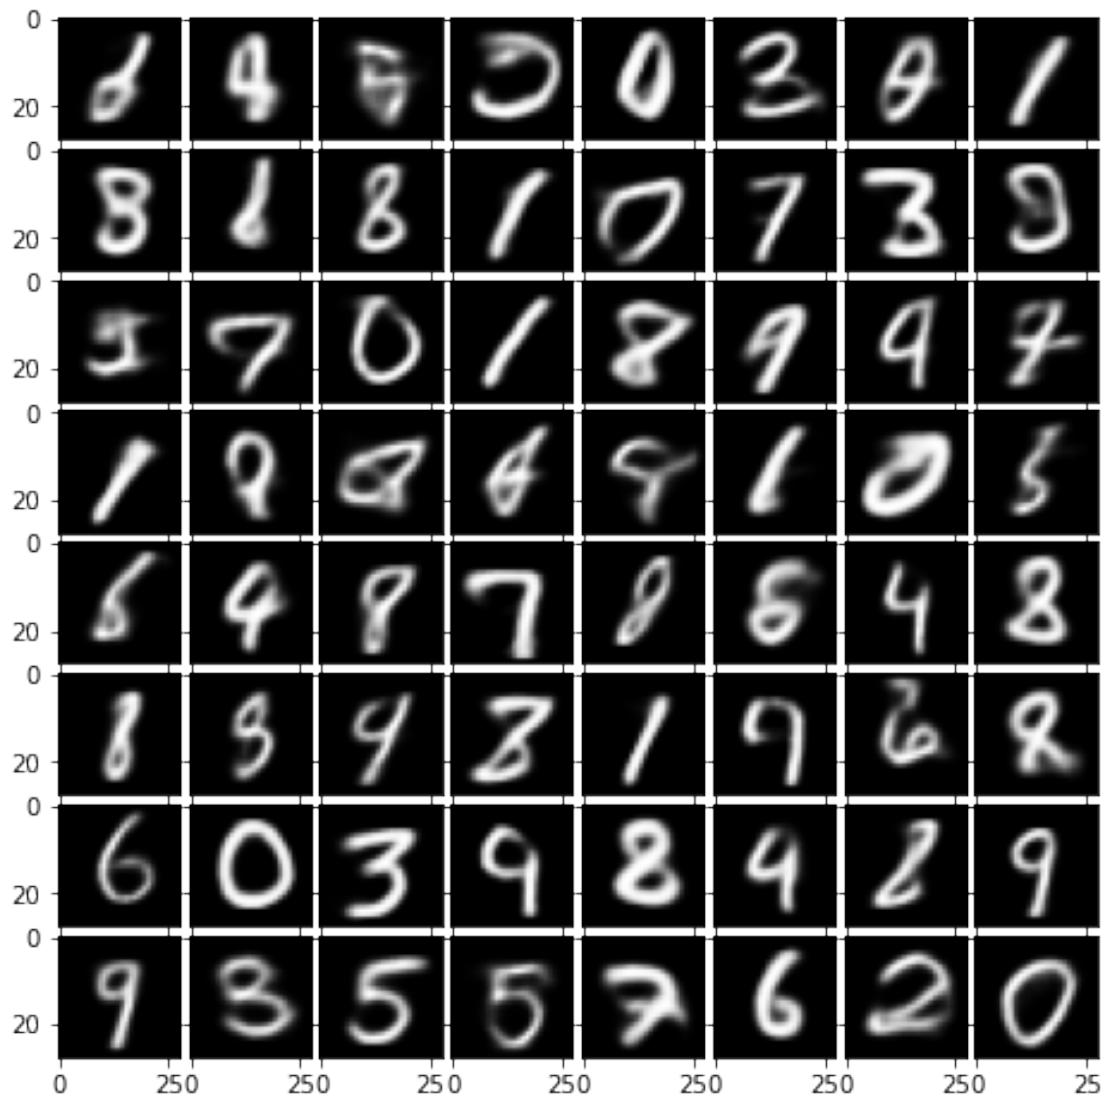


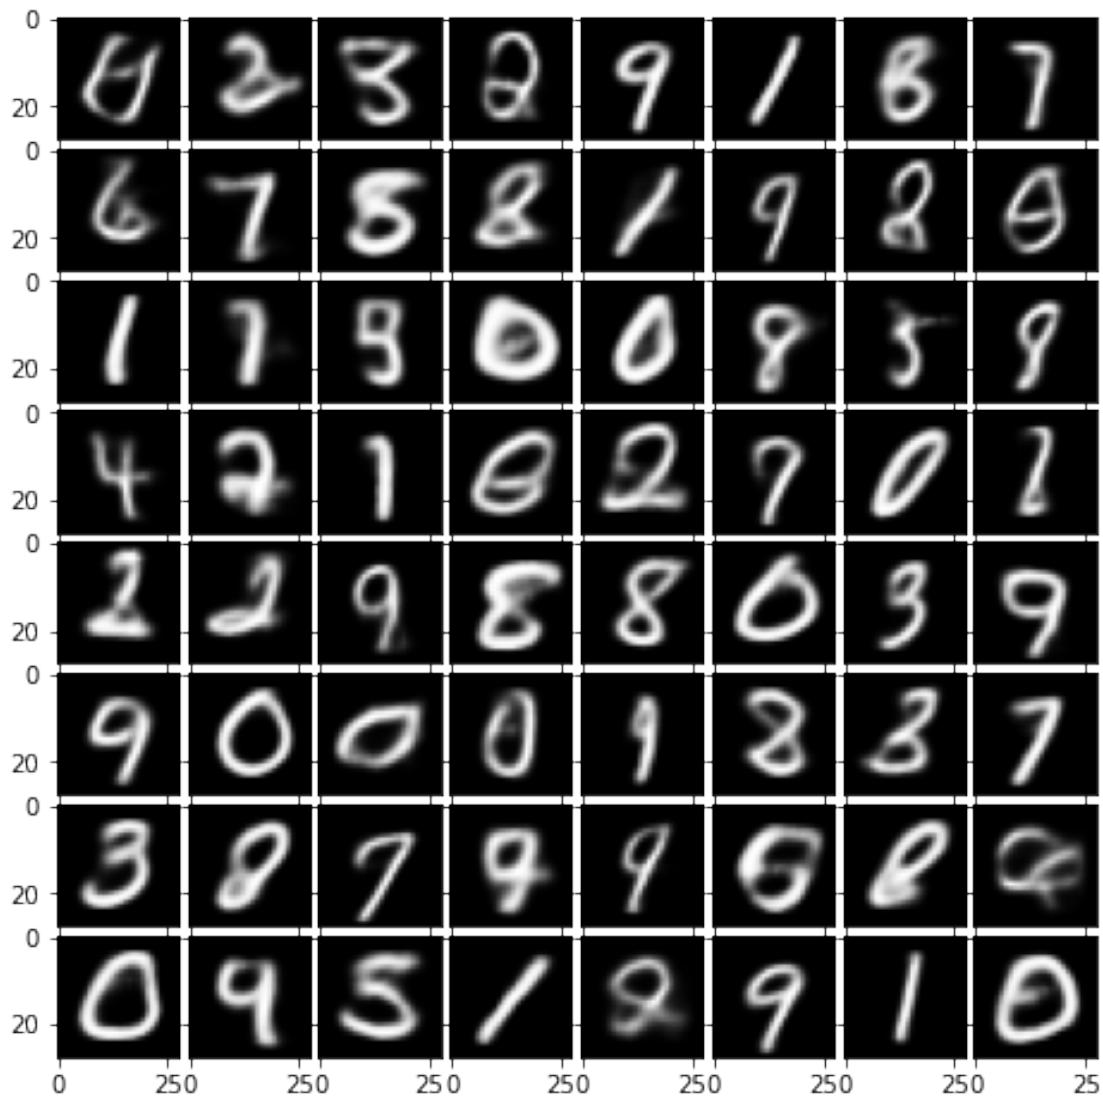












```
[ ]: imgs = model.sample(64).cpu().detach().reshape(64,28,28)

[ ]: import matplotlib.pyplot as plt
from mpl_toolkits.axes_grid1 import ImageGrid

[22]: plt.gray()
fig = plt.figure(figsize=(8., 8.))
grid = ImageGrid(fig, 111, # similar to subplot(111)
                 nrows_ncols=(8, 8), # creates 2x2 grid of axes
                 axes_pad=0.05 # pad between axes in inch.
                 )

for ax, im in zip(grid, imgs):
```

```
# Iterating over the grid returns the Axes.  
ax.imshow(im)  
  
plt.show()
```

<Figure size 432x288 with 0 Axes>

