

NVAE_on_MSCOCO

October 12, 2021

```
[ ]: !wget http://images.cocodataset.org/zips/train2014.zip
!unzip /content/train2014.zip -d /content
!rm /content/train2014.zip
!wget http://images.cocodataset.org/zips/test2014.zip
!unzip /content/test2014.zip -d /content
!rm /content/test2014.zip
!wget http://images.cocodataset.org/zips/val2014.zip
!unzip /content/val2014.zip -d /content
!rm /content/val2014.zip
```

```
[7]: import os
import torch
from PIL import Image
from os import listdir
from sklearn import preprocessing
```

```
[9]: class CocoDataLoader(object):

    def __init__(self, data_dir, transform=None):
        self.transform = transform
        self.image_names = [os.path.join(data_dir, img) for img in
→listdir(data_dir) if os.path.join(data_dir, img)]

    def __len__(self):
        return len(self.image_names)

    def __getitem__(self, idx):

        image = Image.open(self.image_names[idx])

        if self.transform:
            image = self.transform(image)

        return image
```

```
[12]: coco_train_data = '/content/train2014'
coco_valid_data = '/content/val2014'
```

```
coco_test_data = '/content/test2014'
```

```
[10]: import os
import time
import numpy
import json
import torch
import logging
import pandas as pd
import matplotlib.pyplot as plt
from torch import nn
from torchvision.utils import make_grid
from torchvision import transforms
from torch.autograd import Variable
from torch.utils.data import DataLoader, ConcatDataset
from torch.utils.tensorboard import SummaryWriter
from torch.optim.lr_scheduler import ExponentialLR
```

```
[14]: selectedDevice = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Used device: {selectedDevice}")
```

Used device: cuda

```
[85]: image_crop = 375
image_size = 64
mean = [0.5, 0.5, 0.5]
std = [0.5, 0.5, 0.5]
batch_size = 256
learning_rate = 0.001
```

```
[86]: training_data = CocoDataloader(coco_train_data,
                                   transform=transforms.Compose([transforms.
↳ CenterCrop((image_crop, image_crop)),
                                                                    transforms.
↳ Resize((image_size, image_size)),
                                                                    transforms.
↳ RandomHorizontalFlip(),
                                                                    transforms.
↳ ToTensor(),
                                                                    ↵
↳ GreyToColor(image_size),
                                                                    transforms.
↳ Normalize(mean, std)
                                                                    ]))
validation_data = CocoDataloader(coco_valid_data,
                                 transform=transforms.Compose([transforms.
↳ CenterCrop((image_crop, image_crop)),
```

```

transformations.Compose([
    transforms.Resize((image_size, image_size)),
    transforms.ToTensor(),
    transforms.GreyToColor(image_size),
    transforms.Normalize(mean, std)
]))

test_data = CocoDataloader(coco_test_data,
                           transform=transformations.Compose([
    transformations.CenterCrop((image_crop, image_crop)),
    transformations.Resize((image_size, image_size)),
    transformations.RandomHorizontalFlip(),
    transformations.ToTensor(),
    transformations.GreyToColor(image_size),
    transformations.Normalize(mean, std)
]))

train_test_data = ConcatDataset([training_data, test_data])

dataloader_train = DataLoader(train_test_data, batch_size=batch_size,
                               shuffle=True, num_workers=2)
dataloader_valid = DataLoader(validation_data, batch_size=batch_size,
                               shuffle=False, num_workers=2)

```

```

[23]: class depthwise_separable_conv(nn.Module):
        def __init__(self, nin, kernels_per_layer, nout):
            super(depthwise_separable_conv, self).__init__()
            self.depthwise = nn.Conv2d(nin, nin * kernels_per_layer, kernel_size=5,
            padding=2, groups=nin)
            self.pointwise = nn.Conv2d(nin * kernels_per_layer, nout, kernel_size=1)

        def forward(self, x):
            out = self.depthwise(x)
            out = self.pointwise(out)
            return out

```

```

[24]: def swish(x):
        return x * torch.sigmoid(x)

```

```
[25]: class ChannelSELayer(nn.Module):
    def __init__(self, num_channels, reduction_ratio=2):
        super(ChannelSELayer, self).__init__()
        num_channels_reduced = num_channels // reduction_ratio
        self.reduction_ratio = reduction_ratio
        self.fc1 = nn.Linear(num_channels, num_channels_reduced, bias=True)
        self.fc2 = nn.Linear(num_channels_reduced, num_channels, bias=True)
        self.relu = nn.ReLU()
        self.sigmoid = nn.Sigmoid()

    def forward(self, input_tensor):
        batch_size, num_channels, H, W = input_tensor.size()
        squeeze_tensor = input_tensor.view(batch_size, num_channels, -1).
        ↪mean(dim=2)

        fc_out_1 = self.relu(self.fc1(squeeze_tensor))
        fc_out_2 = self.sigmoid(self.fc2(fc_out_1))

        a, b = squeeze_tensor.size()
        output_tensor = torch.mul(input_tensor, fc_out_2.view(a, b, 1, 1))
        return output_tensor
```

```
[26]: class dec_res(nn.Module):
    def __init__(self, in_channel):
        super(dec_res, self).__init__()
        self.bn1 = nn.BatchNorm2d(in_channel)
        self.c1 = nn.
        ↪Conv2d(in_channels=in_channel, out_channels=2*in_channel, kernel_size=1, stride=1, padding=0)
        self.bn2 = nn.BatchNorm2d(2*in_channel)
        self.dc1 =
        ↪depthwise_separable_conv(nin=2*in_channel, kernels_per_layer=3, nout=2*in_channel)
        self.bn3 = nn.BatchNorm2d(2*in_channel)
        self.c2 = nn.
        ↪Conv2d(in_channels=2*in_channel, out_channels=in_channel, kernel_size=1, stride=1, padding=0)
        self.bn4 = nn.BatchNorm2d(in_channel)
        self.SE = ChannelSELayer(in_channel)
    def forward(self, x1):
        x = self.c1(self.bn1(x1))
        x = swish(self.bn2(x))
        x = self.dc1(x)
        x = swish(self.bn3(x))
        x = self.bn4(self.c2(x))
        x = self.SE(x)
        return x+x1
```

```
[27]: class enc_res(nn.Module):
    def __init__(self, in_channel):
        super(enc_res, self).__init__()
        self.bn1 = nn.BatchNorm2d(in_channel)
        self.c1 = nn.
        ↪Conv2d(in_channels=in_channel, out_channels=2*in_channel, kernel_size=3, stride=1, padding=1)
        self.bn2 = nn.BatchNorm2d(2*in_channel)
        self.c2 = nn.
        ↪Conv2d(in_channels=2*in_channel, out_channels=in_channel, kernel_size=3, stride=1, padding=1)
        self.bn3 = nn.BatchNorm2d(in_channel)
        self.SE = ChannelSELayer(in_channel)
    def forward(self, x1):
        x = self.c1(swish(self.bn1(x1)))
        x = self.c2(swish(self.bn2(x)))
        x = self.SE(x)
        return x+x1
```

```
[79]: class NVAE(nn.Module):
    def __init__(self, start_channel, original_dim):
        super(NVAE, self).__init__()
        self.original_dim = original_dim
        self.conv1 = nn.
        ↪Conv2d(in_channels=start_channel, out_channels=8, kernel_size=3, stride=1, padding=1)
        self.encblock1 = enc_res(8)
        self.dsconv1 = nn.
        ↪Conv2d(in_channels=8, out_channels=8, kernel_size=2, stride=2, padding=0)
        self.encblock2 = enc_res(8)
        self.dsconv2 = nn.
        ↪Conv2d(in_channels=8, out_channels=8, kernel_size=2, stride=2, padding=0)

        self.qmu1 = nn.
        ↪Linear(original_dim*original_dim*2, original_dim*original_dim*2)
        self.qvar1 = nn.
        ↪Linear(original_dim*original_dim*2, original_dim*original_dim*2)

        self.qmu0 = nn.Linear(original_dim*original_dim//
        ↪2, original_dim*original_dim//2)
        self.qvar0 = nn.Linear(original_dim*original_dim//
        ↪2, original_dim*original_dim//2)

        self.pmu1 = nn.
        ↪Linear(original_dim*original_dim*2, original_dim*original_dim*2)
        self.pvar1 = nn.
        ↪Linear(original_dim*original_dim*2, original_dim*original_dim*2)

        self.decblock1 = dec_res(8)
```

```

        self.usconv1 = nn.
        ↪ConvTranspose2d(in_channels=8,out_channels=8,kernel_size=2,stride=2,padding=0)
        self.decblock2 = dec_res(16)
        self.usconv2 = nn.
        ↪ConvTranspose2d(in_channels=16,out_channels=16,kernel_size=2,stride=2,padding=0)
        self.decblock3 = dec_res(16)
        self.finconv = nn.
        ↪Conv2d(in_channels=16,out_channels=start_channel,kernel_size=3,stride=1,padding=1)

    def forward(self,x):
        z1 = self.dsconv1(self.encblock1(self.conv1(x)))
        z0 = self.dsconv2(self.encblock2(z1))

        qmu0 = self.qmu0(z0.reshape(z0.shape[0],self.original_dim*self.original_dim/
        ↪/2))
        qvar0 = self.qvar0(z0.reshape(z0.shape[0],self.original_dim*self.
        ↪original_dim//2))

        qmu1 = self.qmu1(z1.reshape(z1.shape[0],self.original_dim*self.
        ↪original_dim*2))
        qvar1 = self.qvar1(z1.reshape(z1.shape[0],self.original_dim*self.
        ↪original_dim*2))

        stdvar0 = qvar0.mul(0.5).exp_()
        stdvar1 = qvar1.mul(0.5).exp_()

        e0 = torch.randn(qmu0.shape).to(device)
        ez0 = qmu0+e0*stdvar0
        ez0 = ez0.reshape(ez0.shape[0],8,self.original_dim//4,self.original_dim//4)
        ez1 = self.usconv1(self.decblock1(ez0))

        pmu1 = self.pmu1(ez1.reshape(ez1.shape[0],self.original_dim*self.
        ↪original_dim*2))
        pvar1 = self.pvar1(ez1.reshape(ez1.shape[0],self.original_dim*self.
        ↪original_dim*2))

        pstdvar1 = pvar1.mul(0.5).exp_()

        e2 = torch.randn(qmu1.shape).to(device)
        ez2 = pmu1+qmu1 + e2*pstdvar1*stdvar1
        ez2 = ez2.reshape(ez2.shape[0],8,self.original_dim//2,self.original_dim//2)

        final = torch.cat((ez1,ez2),1)

        recons = nn.Sigmoid()(self.finconv(self.decblock3(self.usconv2(self.
        ↪decblock2(final)))))

```

```

        return qmu0,qvar0,qmu1,qvar1,pmu1,pvar1,recons

    def sample(self,bs):
        e = torch.randn([bs,8,self.original_dim//4,self.original_dim//4]).to(device)
        ez1 = self.usconv1(self.decblock1(e))

        pmu1 = self.pmu1(ez1.reshape(ez1.shape[0],self.original_dim*self.
↪original_dim*2))
        pvar1 = self.pvar1(ez1.reshape(ez1.shape[0],self.original_dim*self.
↪original_dim*2))

        stdvar1 = pvar1.mul(0.5).exp_()

        e1 = torch.randn([ez1.shape[0],self.original_dim*self.original_dim*2]).
↪to(device)
        e1 = pmu1 + e1*stdvar1
        e1 = e1.reshape(e1.shape[0],8,self.original_dim//2,self.original_dim//2)
        recons = nn.Sigmoid()(self.fconv(self.decblock3(self.usconv2(self.
↪decblock2(torch.cat((ez1,e1),1))))))

        return recons

    def loss(self,x):
        qmu0,qvar0,qmu1,qvar1,pmu1,pvar1,recons = self.forward(x)
        klz0 = 0.5*torch.sum(torch.square(qmu0)+qvar0.exp()-qvar0-1)/x.shape[0]
        klz1 = 0.5*torch.sum(torch.square(qmu1)/pvar1.exp()+qvar1.exp()-qvar1-1)
        reconsloss = nn.BCELoss()(recons,x)
        return klz0,klz1,reconsloss

```

```
[89]: model = NVAE(3,64).to(selectedDevice)
```

```
[90]: optim = torch.optim.Adamax(model.parameters())
device = selectedDevice
```

```
[91]: epochs=20
```

```
[83]: import matplotlib.pyplot as plt
from mpl_toolkits.axes_grid1 import ImageGrid
```

```
[92]: for epoch in range(epochs):
        minloss = 1
        running_kl0_loss=0
        running_recons_loss=0
        running_kl1_loss=0
        num_images=0
        for i, img in enumerate(dataloader_train):
```

```

img = img.to(selectedDevice)
optim.zero_grad()
klz0,klz1,recons = model.loss(img)
loss=recons+epoch*0.0001*klz0+epoch*0.0001*klz1
loss.backward()
optim.step()
running_kl0_loss = running_kl0_loss + klz0.item()*len(img)
running_kl1_loss = running_kl1_loss + klz1.item()*len(img)
running_recons_loss = running_recons_loss + recons.item()*len(img)

num_images= num_images+len(img)
print('epoch: '+str(epoch)+' kl0_loss: '+str(running_kl0_loss/num_images)+'
↳recons_loss: '+str(running_recons_loss/num_images)+' kl1_loss:
↳'+str(running_kl1_loss/num_images))
imgs = model.sample(64).cpu().detach().reshape(64,28,28)
plt.gray()
fig = plt.figure(figsize=(8., 8.))
grid = ImageGrid(fig, 111, nrows_ncols=(8, 8), axes_pad=0.05)

for ax, im in zip(grid, imgs):
    ax.imshow(im)

```

```

-----
RuntimeError                                Traceback (most recent call last)
<ipython-input-92-a501f81f0bd5> in <module>()
    11     loss=recons+epoch*0.0001*klz0+epoch*0.0001*klz1
    12     loss.backward()
----> 13     optim.step()
    14     running_kl0_loss = running_kl0_loss + klz0.item()*len(img)
    15     running_kl1_loss = running_kl1_loss + klz1.item()*len(img)

/usr/local/lib/python3.7/dist-packages/torch/optim/optimizer.py in
↳wrapper(*args, **kwargs)
    86         profile_name = "Optimizer.step#{}.step".format(obj.
↳__class__.__name__)
    87         with torch.autograd.profiler.
↳record_function(profile_name):
----> 88             return func(*args, **kwargs)
    89         return wrapper
    90

/usr/local/lib/python3.7/dist-packages/torch/autograd/grad_mode.py in
↳decorate_context(*args, **kwargs)
    26     def decorate_context(*args, **kwargs):
    27         with self.__class__():
----> 28             return func(*args, **kwargs)
    29         return cast(F, decorate_context)

```


30

```
/usr/local/lib/python3.7/dist-packages/torch/optim/adamax.py in step(self,
↳ closure)
    94         beta2=beta2,
    95         lr=lr,
--> 96         weight_decay=weight_decay)

    97
    98     return loss

/usr/local/lib/python3.7/dist-packages/torch/optim/_functional.py in
↳ adamax(params, grads, exp_avgs, exp_infs, state_steps, eps, beta1, beta2, lr,
↳ weight_decay)
    315     norm_buf = torch.cat([
    316         exp_inf.mul_(beta2).unsqueeze(0),
--> 317         grad.abs().add_(eps).unsqueeze_(0)
    318     ], 0)
    319     torch.amax(norm_buf, 0, keepdim=False, out=exp_inf)
```

RuntimeError: CUDA error: device-side assert triggered
CUDA kernel errors might be asynchronously reported at some other API call,so
↳ the stacktrace below might be incorrect.
For debugging consider passing CUDA_LAUNCH_BLOCKING=1.

[68]: image.size

[68]: (640, 425)

[]: